



# **OPERATING SYSTEMS**

## **Week 3**

**Much of the material on these slides comes from the recommended textbook by William Stallings**

# Detailed content

## Weekly program

- ✓ Week 1 – Operating System Overview
- ✓ Week 2 – Processes and Threads



### ☐ Week 3 – Scheduling

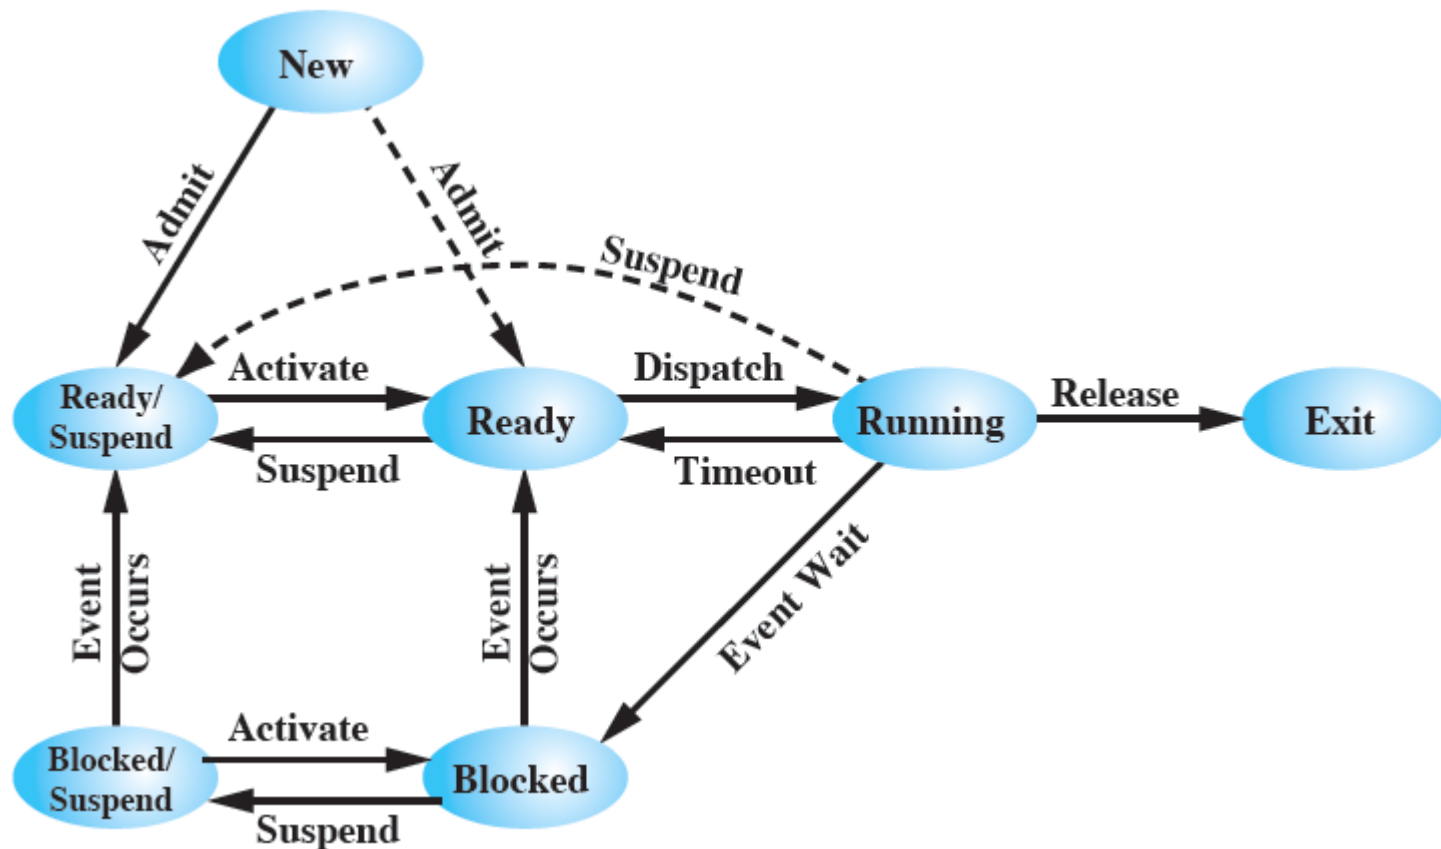
- ☐ Week 4 – Real-time System Scheduling and Multiprocessor Scheduling
- ☐ Week 5 – Concurrency: Mutual Exclusion and Synchronization
- ☐ Week 6 – Concurrency: Deadlock and Starvation
- ☐ Week 7 – Memory Management
- ☐ Week 8 – Disk and I/O Scheduling
- ☐ Week 9 – File Management
- ☐ Week 10 – Real-world Operating Systems: Embedded and Security
- ☐ Week 11 – Real-world Operating Systems
- ☐ Week 12 – Revision of the course
- ☐ Week 13 – Extra revision (if needed)

# Key Concepts From Last Week

- A process is an entity consisting of two essential elements: program code and a set of data
- Processes Control Block is a data structure that stores all relevant information related to a process
- A process passes through different states throughout its lifetime
  - Two-state process model
  - Five-state model
  - Seven-state model

# Key Concepts From Last Week

4



# Key Concepts from last week

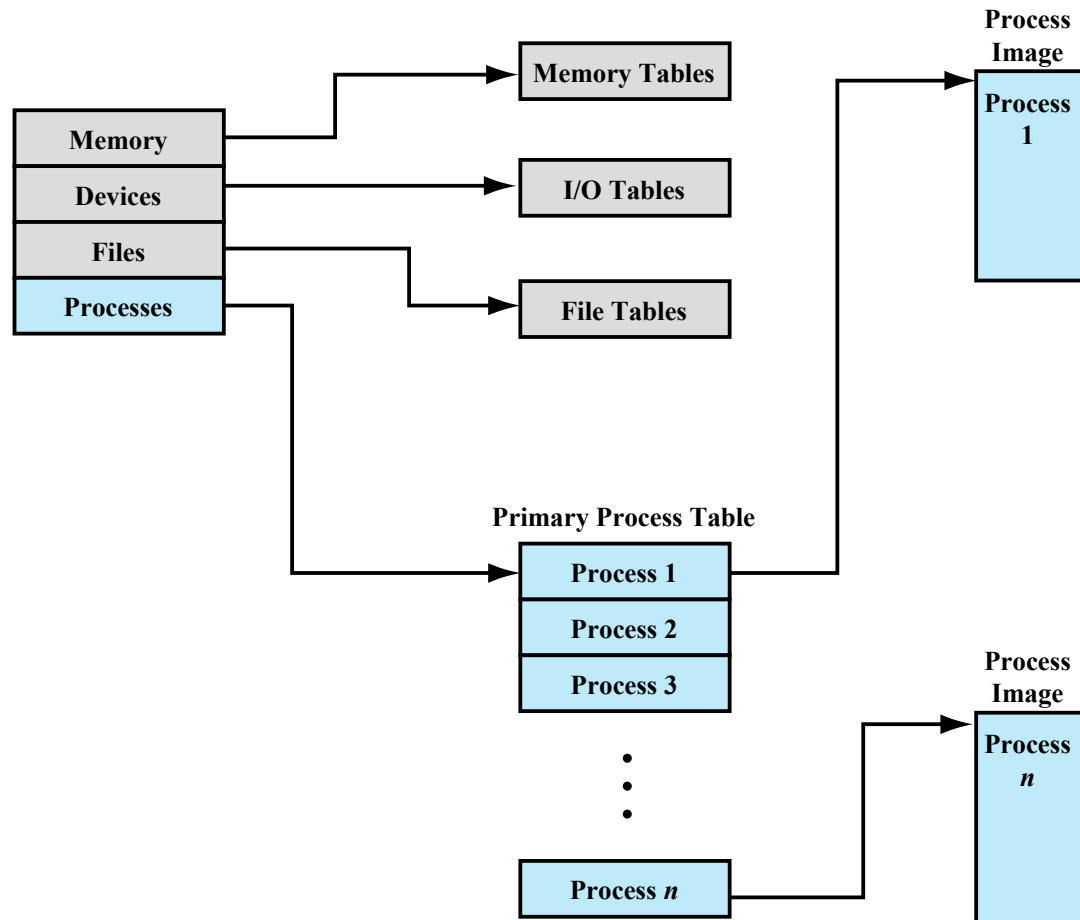


Figure 3.11 General Structure of Operating System Control Tables

# Key Concepts from last week

6

- Significant amount of data is stored in PCB
  - Process ID, Process State Information, Process Control Information
- The difference between context switching and mode switching
- In multithreading environment: Process is related to resource ownership and Thread (light weight process) is related to program execution
- Types of threads: ULT, KLT, Combined

# Week 03 Lecture Outline

## Scheduling

- ❑ Scheduling – what, why and when?
- ❑ Types of Scheduling
- ❑ Short time Scheduling/CPU Scheduling
- ❑ Criteria for CPU Scheduling
- ❑ Use of priority
- ❑ Scheduling:
  - ❑ Selection function and
  - ❑ Decision mode
- ❑ Scheduling Algorithms
  - ❑ First Come First Served (FCFS)
  - ❑ Round Robin (RR)
  - ❑ Shortest Process Next (SPN)
  - ❑ Shortest Remaining Time (SRT)
  - ❑ Highest Response Ration Next (HRRN)
  - ❑ Feedback (FB)



Videos to watch before lecture



# What is Scheduling?

- On a multi-programmed system
  - We may have more than one *Ready* process
- On a batch system
  - We may have many jobs waiting to be run
- On a multi-user system
  - We may have many users concurrently using the system
- The *scheduler* decides who to run next.
  - The process of choosing is called **scheduling**.





# Is scheduling important?

- It is not in certain scenarios
  - If you have no choice
- Early systems
  - Usually batching
  - Scheduling algorithm simple
    - Run next on tape or next on punch tape
  - Only one thing to run
- Simple PCs
  - Only ran a word processor, etc....
- Simple Embedded Systems
  - TV remote control, washing machine, etc....




# Is scheduling important?

- It is in more complex scenarios
  - Multitasking/Multi-user systems
- Example
  - Email daemon takes 2 seconds to process an email
  - User clicks button on application.
    - **Scenario 1:** Run daemon, then application
      - System appears really sluggish to the user
    - **Scenario 2:** Run application, then daemon
      - Application appears really responsive, small email delay is unnoticed
- Scheduling decisions can have a dramatic effect on the perceived performance of the system
  - Can also affect correctness of a system with deadlines



# Terminology



## a) CPU-Bound process

- Spends most of its computing
- Time to completion largely determined by received CPU time

## b) I/O-Bound process

- Spend most of its time waiting for I/O to complete
  - Small bursts of CPU to process I/O and request next I/O
- Time to completion largely determined by I/O request time



# Observations

1. Generally, technology is increasing CPU speed much faster than I/O speed
  - CPU bursts becoming shorter, I/O waiting is relatively constant
  - Processes are becoming more I/O bound
2. We need a mix of CPU-bound and I/O-bound processes to keep both CPU and I/O systems busy
3. Process can go from CPU- to I/O-bound (or vice versa) in different phases of execution



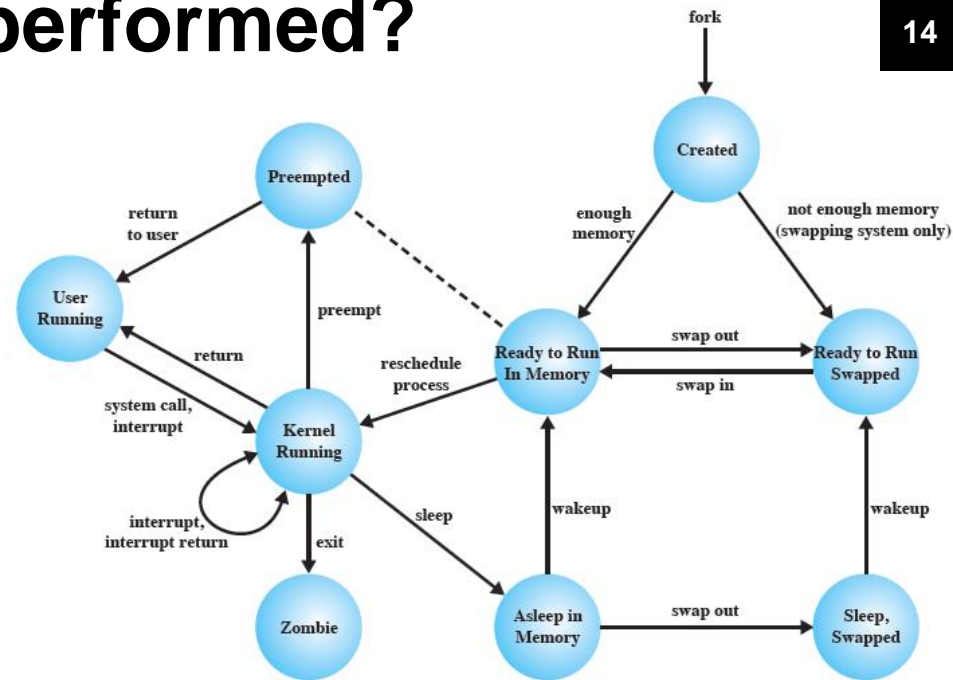
# Observations

4. Choosing to run an I/O-bound process delays a CPU-bound process by very little
5. Choosing to run a CPU-bound process prior to an I/O-bound process delays the next I/O request significantly
  - No overlap of I/O waiting with computation
  - Results in device (disk) not as busy as possible
6. Generally, favour I/O-bound processes over CPU-bound processes



# When is scheduling performed?

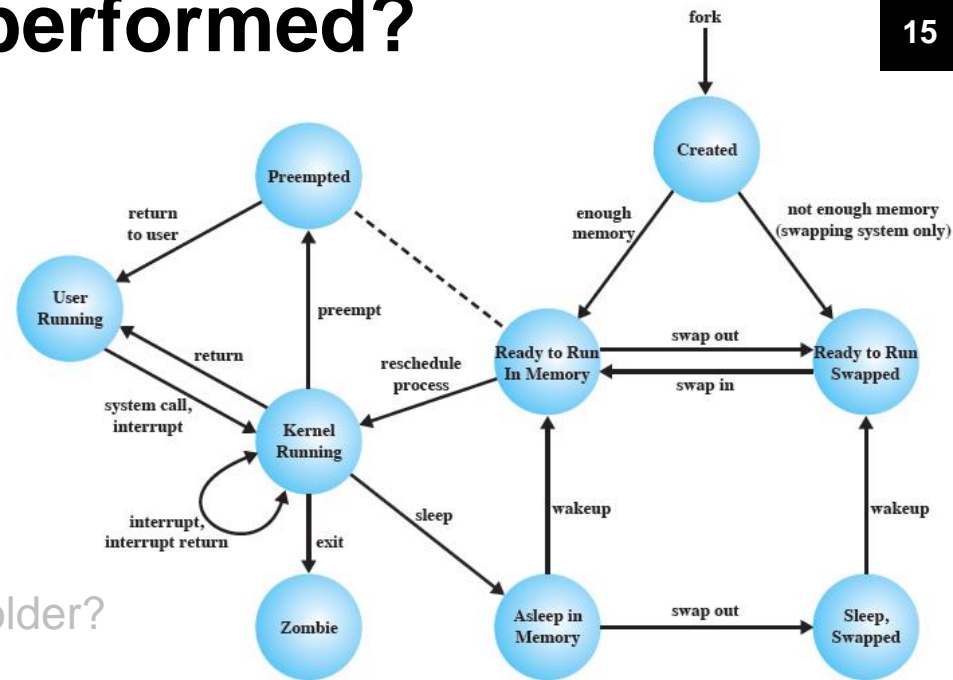
- A new process
  - Run the parent or the child?
- A process exits
  - Who runs next?
- A process waits for I/O
  - Who runs next?
- A process blocks on a lock
  - Who runs next? The lock holder?
- An I/O interrupt occurs
  - Who do we resume, the interrupted process or the process that was waiting?
- On a timer interrupt? (Preemptive versus Non-preemptive Scheduling)





# When is scheduling performed?

- A new process
  - Run the parent or the child?
- A process exits
  - Who runs next?
- A process waits for I/O
  - Who runs next?
- A process blocks on a lock
  - Who runs next? The lock holder?
- An I/O interrupt occurs
  - Who do we resume, the interrupted process or the process that was waiting?
- On a timer interrupt? (Preemptive versus Non-preemptive Scheduling)



- Generally, a scheduling decision is required when a process (or thread) can no longer continue, or when an activity results in more than one ready process.

# Types of scheduling

There are four kinds of scheduling in an OS:

1. **Long term scheduling**

- The decision to add to the pool of processes to be executed

2. **Medium term scheduling**

- The decision to add to the number of processes that are partially or fully in main memory

3. **Short term scheduling**

- The decision as to which available process will be executed by the processor


4. **I/O scheduling**

- The decision as to which process's pending I/O request shall be handled by an available I/O device.



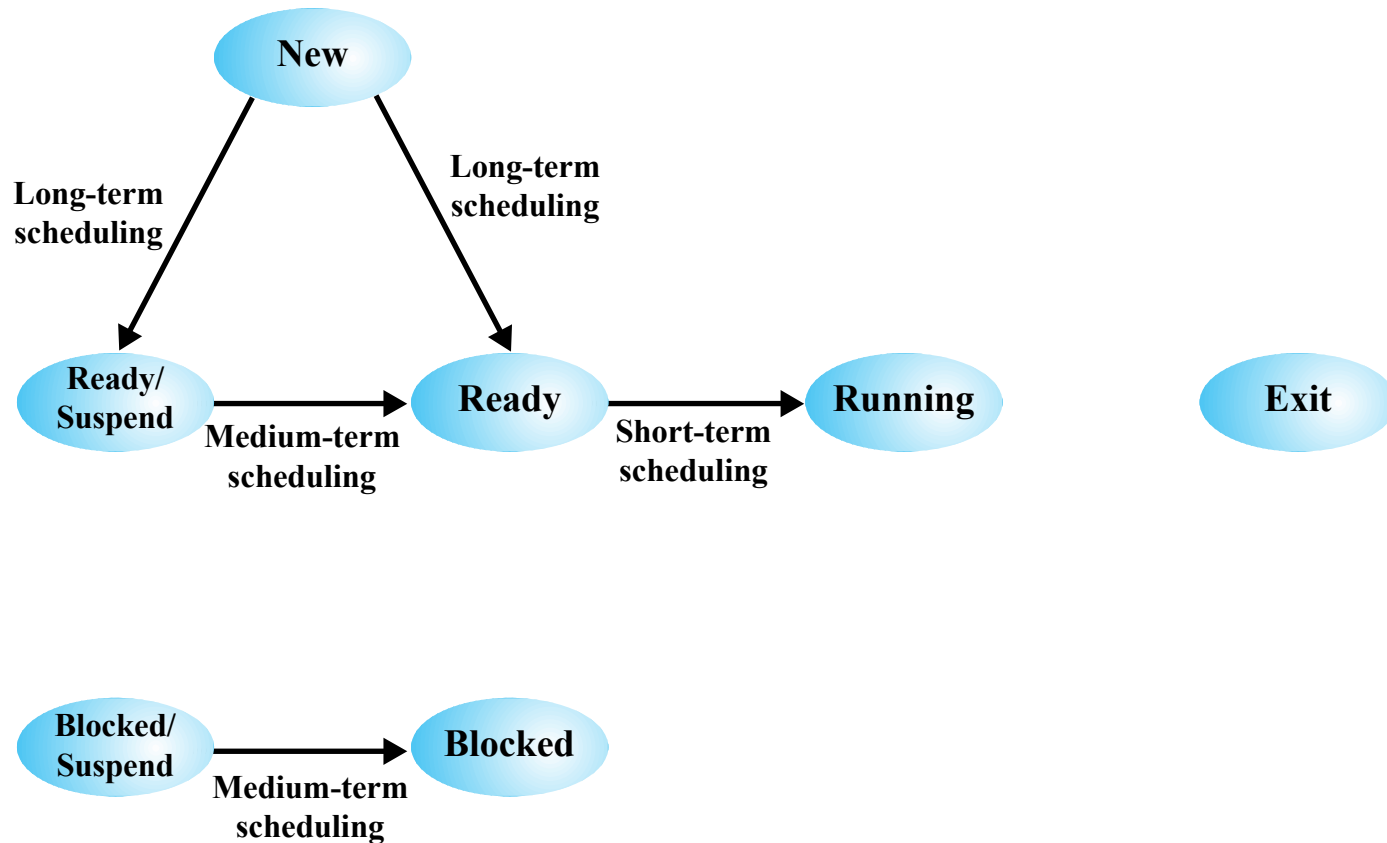
# Processor Scheduling

- Aim is to assign processes to be executed by the processor in a way that meets system objectives, such as response time, throughput, and processor efficiency
- Broken down into three separate functions:



- The names suggest the relative time scales with which these functions are performed

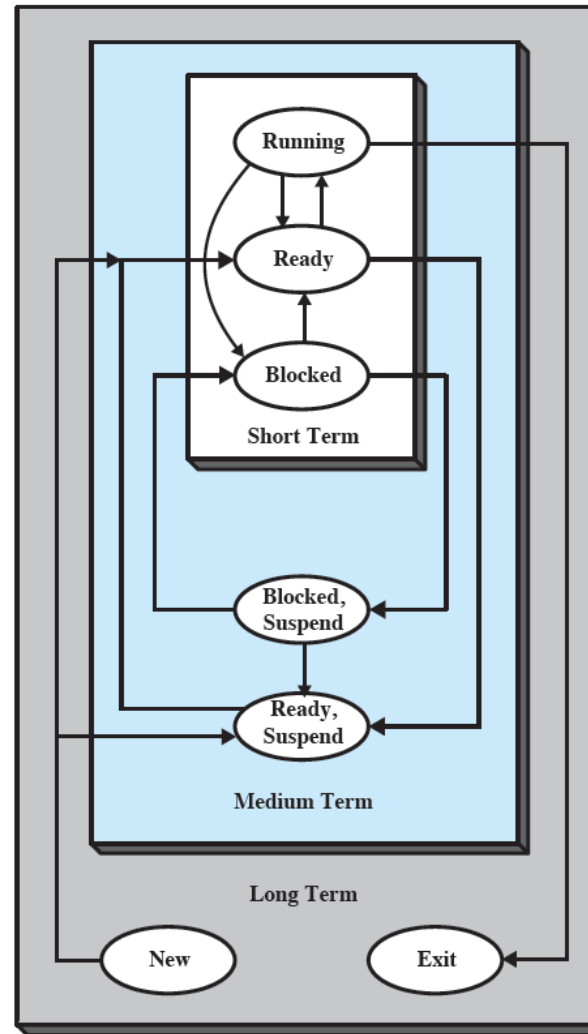
# Processor Scheduling



**Figure 9.1 Scheduling and Process State Transitions**

# Processor Scheduling

19



# Processor Scheduling

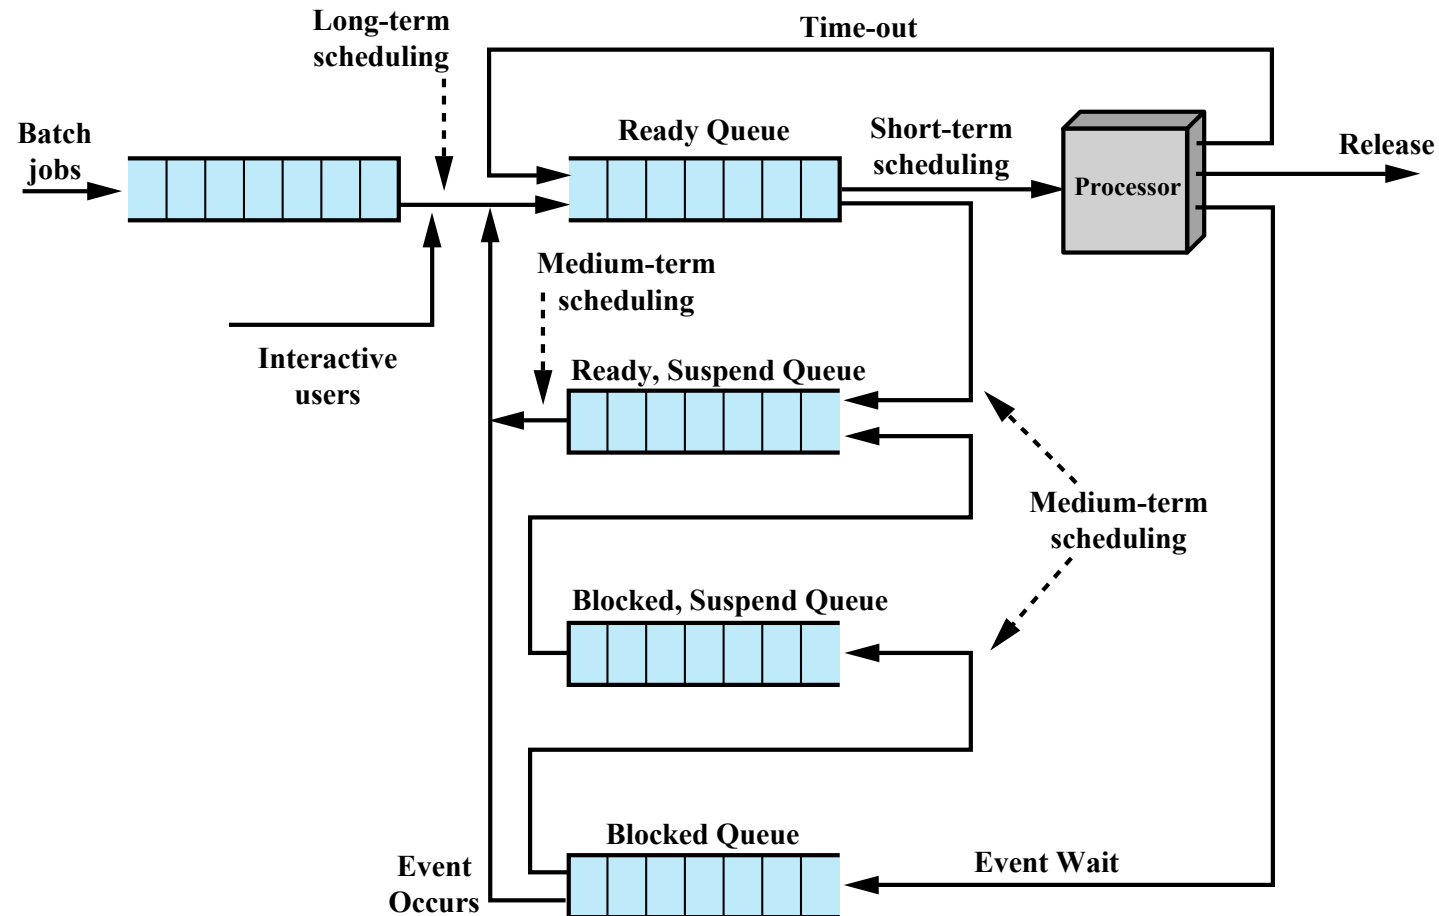
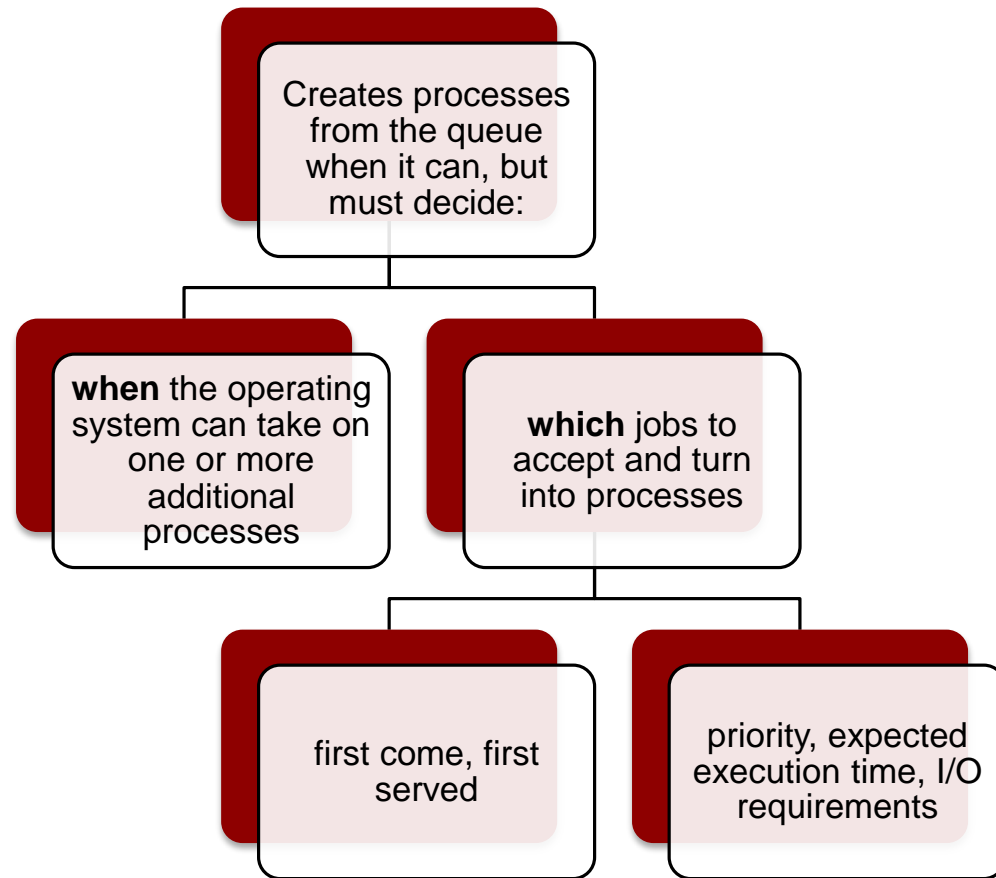


Figure 9.3 Queuing Diagram for Scheduling

# Scheduling: Long-Term

- Determines which programs are admitted to the system for processing
- Controls the degree of multiprogramming
  - the more processes that are created, the smaller the percentage of time that each process can be executed
  - may limit to provide satisfactory service to the current set of processes




# Scheduling: Long-Term

- For interactive-programs in a time-sharing system
  - Process creation request generated when user attempt to connect to a system
  - Are not queued up and kept waiting until the system is saturated
    - According to criteria
  - If saturated then refused with an error message

# Scheduling: Medium-Term

- Part of the swapping function
- Swapping-in decisions are based on the need to manage the degree of multiprogramming
  - considers the memory requirements of the swapped-out processes
- Executes more frequently than long term scheduler.
- It is part of the **memory management** system.
  - We will look at this topic later.



# Scheduling: Short-Term

- Known as the dispatcher
- Executes most frequently
- Makes the fine-grained decision of which process to execute next
- Invoked when an event occurs that may lead to the blocking of the current process or that may provide an opportunity to preempt a currently running process in favor of another
  - Examples:
    - Clock interrupts
    - I/O interrupts
    - Operating system calls
    - Signals (e.g., semaphores)



# Short Term Scheduling Criteria

- Main objective is to allocate processor time to optimize certain aspects of system behavior
- A set of criteria is needed to evaluate the scheduling policy
- **User-oriented criteria**
  - relate to the behaviour of the system as perceived by the individual user or process (such as response time in an interactive system)
  - important on virtually all systems
  - **Focus:** Service provided to the user
- **System-oriented criteria**
  - focus in on effective and efficient utilization of the processor (rate at which processes are completed)
  - generally of minor importance on single-user systems
  - **Focus:** System performance



# Scheduling Policies

There are many kinds of short term scheduling algorithms.

Schedulers have two main parameters:

- **Selection function:**
- **Decision mode:**

# Selection function

- Determines which process, among ready processes, is selected next for execution
- May be based on **priority**, **resource requirements**, or the **execution characteristics** of the process
- If based on execution characteristics then important quantities are:
  - $w$  = time spent in system so far, waiting
  - $e$  = time spent in execution so far
  - $s$  = total service time required by the process, including  $e$ ;  
generally, this quantity must be estimated or supplied by the user



# Decision mode

- Specifies the instants in time at which the selection function is exercised
- Two categories:
  - Nonpreemptive
  - Preemptive




# Nonpreemptive vs. Preemptive

- **Nonpreemptive**

- Once a process is in the running state, it will continue until it terminates or blocks itself (e.g. for I/O)

- **Preemptive**

- Currently running process may be interrupted and moved to ready state by the OS
- Preemption may occur when new process arrives, on an interrupt, or periodically



# Nonpreemptive vs. Preemptive

- **Nonpreemptive**
  - Pros:
    - Less overhead (less switching)
  - Cons:
    - May monopolize
- **Preemptive**
  - Pros:
    - Does not monopolize
    - Better overall service to all processes
  - Cons:
    - More overhead (more switching)
      - Overhead minimization using efficient switching



# Scheduling Algorithms

- **First Come First Served (FCFS)**
- **Round Robin (RR)**
- **Shortest Process Next (SPN)**
- **Shortest Remaining Time (SRT)**
- **Highest Response Ration Next (HRRN)**
- **Priority Scheduling (PS)**
- **Feedback (FB)**
- **Fair Share Scheduling (FSS)**



# Performance Indices

- **Arrival Time ( $T_a$ ):** Time when process enter the system (short time scheduling)
- **Service Time ( $T_s$ ):** Total execution time required by the process
- **Turnaround Time ( $T_r$ ):** Total time that the items spends in the system
- **Normalized Turnaround Time:**  $T_r/T_s$ : Relative delay experienced by a process. Minimum value is 1.0; Increasing values correspond to decreasing level of service.



# Running Example

33

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2




# First-Come-First-Served (FCFS)

- Simplest scheduling policy
- Also known as first-in-first-out (FIFO) or a strict queuing scheme
- When the current process ceases to execute, the process that has been in the Ready queue the longest is selected
- Non-preemptive policy
- It may be implemented by using the ready queue as a FIFO queue.
- Performs much better for long processes than short ones
- Tends to favor processor-bound processes over I/O-bound processes



# FCFS - Example

FCFS on this set of jobs:




Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_S$ )	3	6	4	5	2	Mean
FCFS						
Finish Time	3	9	13	18	20	
Turnaround Time ( $T_T$ )	3	7	9	12	12	8.60
$T_T/T_S$	1.00	1.17	2.25	2.40	6.00	2.56

# FCFS Performance

- FCFS is simple but performs badly if the job mix contains jobs of widely different characteristics, for example:

process	A	B	C	D
arrival time	0	1	2	3
service time	1	100	1	100
start time	0	1	101	102
finish time	1	101	102	202
turnaround time $T_r$	1	100	100	199
$T_r/T_s$	1.0	1.0	100	2.0



- Here the long jobs (B, and D) get a reasonable turnaround time, but the short job (C) has an unreasonable turnaround time.

# FCFS Performance

- FCFS also has the **disadvantage** that a CPU bound job can monopolise the processor, leaving I/O devices idle.
  - Suppose there is one processor bound job and many I/O bound jobs
    - When CPU bound job runs all I/O bound jobs wait and I/O devices are idle
    - When the CPU bound job finishes, the I/O bound jobs quickly become blocked waiting and the CPU becomes idle.
    - May result in inefficient use of both the processor and the I/O devices
    - **Convoy effect**
- FCFS is not an attractive scheduling policy on its own
- FCFS is best used in **combination** with another policy


# Round Robin

- Stop monopolization of length job by preemption
- In **Round Robin scheduling**, (RR) or **time slicing**, a clock regularly interrupts and pre-empts the running process.
- The next process to run is chosen by an FCFS strategy.
- The main parameter is the **length of the time slice**. A short time slice ensures that brief processes move through the system rapidly, but means many interrupts.
- What happens when time slice becomes infinite?



# RR - Example


- Example for time slice of 1 unit:



Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_S$ )	3	6	4	5	2	Mean
<b>RR <math>q = 1</math></b>						
Finish Time	4	18	17	20	15	
Turnaround Time ( $T_T$ )	4	16	13	14	7	10.80
$T_T/T_S$	1.33	2.67	3.25	2.80	3.50	2.71

# RR - Example

- Example with time slice of 4 units:

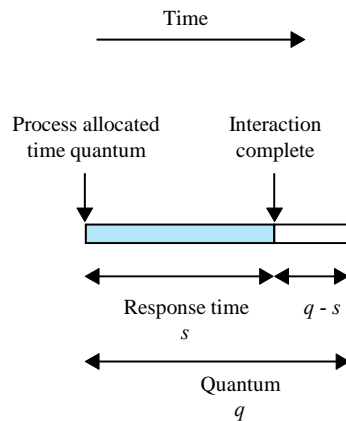


Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_s$ )	3	6	4	5	2	Mean
RR $q = 4$						
Finish Time	3	17	11	20	19	
Turnaround Time ( $T_T$ )	3	15	7	14	11	10.00
$T_T/T_s$	1.00	2.5	1.75	2.80	5.50	2.71

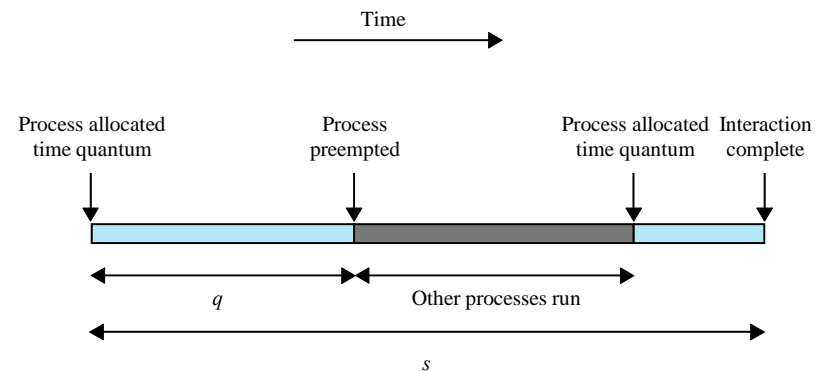


# Effect of preemption time quantum

- The time slice may be chosen to be a little **larger than the time required for a typical interaction**; that is, most processes should require only one time slice.



(a) Time quantum greater than typical interaction




(b) Time quantum less than typical interaction

# Round Robin Performance

- Round Robin is effective for general purpose time-sharing systems.
- Each context switch has the OS using the CPU instead of the user process
  - give up CPU, save all info, reload w/ status of incoming process
  - Say 20 ms quantum length, 5 ms context switch
  - Waste of resources
    - 20% of CPU time (5/25) for context switch
  - If 500 ms quantum, better use of resources
    - 1% of CPU time (5/505) for context switch
    - Bad if lots of users in system – interactive users waiting for CPU
  - Balance found depends on job mix
- Still favors CPU-bound processes
  - An I/O bound process uses the CPU for a time less than the time quantum and then is blocked waiting for I/O
  - A CPU-bound process runs for its whole time slice and goes back into the ready queue (in front of the blocked processes)

# Virtual Round Robin (VRR)




# Virtual Round Robin (VRR)

- **Virtual Round Robin (VRR)** avoids unfairness of RR.
- New processes arrive & join Ready queue - FCFS.
- Timed-out running processes rejoin Ready queue.
- I/O blocked processes join appropriate I/O queue.
  - (So far this is standard)
- NEW FEATURE: add FCFS *Auxiliary* queue.
- Processes released from I/O wait queue join *Auxiliary* queue.
- Scheduler dispatches processes from *Auxiliary* queue before processes from Ready queue. But these processes only run with up to the time remaining from their last basic time quantum (before they were blocked).

# SPN

- The **Shortest Process Next** (SPN) policy always chooses the process which has the shortest expected running time.
- This policy is **non-pre-emptive**.
- The process with the shortest expected processing time is selected next (Shortest Job First – SJF)
- A short process will jump to the head of the queue

# SPN - Example




Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_S$ )	3	6	4	5	2	Mean
SPN						
Finish Time	3	9	15	20	11	
Turnaround Time ( $T_T$ )	3	7	11	14	3	7.60
$T_T/T_S$	1.00	1.17	2.75	2.80	1.50	1.84

- This example assumes that the processor knows the service time before it runs the processes.

# SPN - Performance

- Gives minimum average waiting time for a given set of processes.
- Possibility of starvation for longer processes
- Not suitable for time-sharing
  - No preemption
- The same bad example as for FCFS applies: with a mixture of very long and very short processes, SPN can be **unfair to short processes**.

process	A	B	C	D
arrival time	0	1	2	3
service time	1	100	1	100
start time	0	1	101	102
finish time	1	101	102	202
turnaround time $T_r$	1	100	100	199
$T_r/T_s$	1.0	1.0	100	2.0




# SRT

- The **Shortest Remaining Time** (SRT) scheduler chooses a process with the shortest expected remaining time.
- This is a kind of **preemptive version of SPN**.
- When a new process becomes available, the OS may **preempt** a running process if the expected time to completion for the currently running process is longer than the expected run time of the new process.
- As with SPN:
  - the OS must estimate runtimes
  - long processes may be starved



# SRT Example



Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_s$ )	3	6	4	5	2	Mean
<b>SRT</b>						
Finish Time	3	15	8	20	10	
Turnaround Time ( $T_r$ )	3	13	4	14	2	7.20
$T_r/T_s$	1.00	2.17	1.00	2.80	1.00	1.59

Notice that the shortest 3 processes all receive immediate attention, and hence have a normalised turnaround time of 1.0 .

# SRT - Performance

## Pros:

- Does not have the bias in favour of long processes like in FCFS
- No additional interrupt are generated like in RR
- Turnaround time is better than SPN because short jobs get immediate attention.


## Cons:

- Elapsed times must be recorded; this increases overhead.

# HRRN

- The **Highest Response Ratio Next** (HRRN) policy chooses the next process to run to be the one with the highest value of RR, defined below.
- The policy is **non-pre-emptive**.
- Define  $RR = \frac{w+s}{s}$  where  
 $w$  = time spent waiting, and  
 $s$  = expected service time.
- Thus, at the completion of a job,  $RR = T_r/T_s$ , which is the value we are trying to minimise.

# HRRN - Example



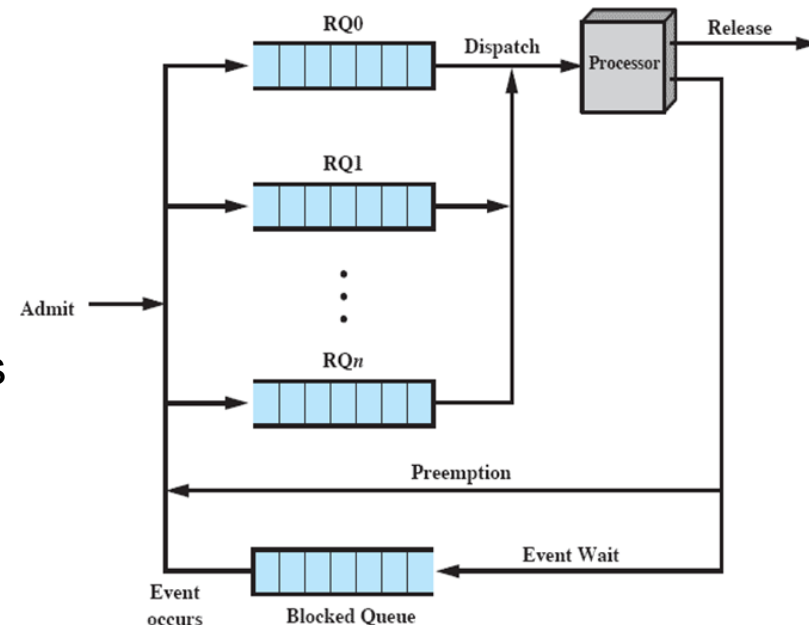
Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_s$ )	3	6	4	5	2	Mean
<b>HRRN</b>						
Finish Time	3	9	13	20	15	
Turnaround Time ( $T_T$ )	3	7	9	14	7	8.00
$T_T/T_s$	1.00	1.17	2.25	2.80	3.5	2.14

# HRRN - Performance

- It considers the age of the process
- If a process is waiting and getting no service, then  $w$  increases while  $s$  remains constant, so  $RR = \frac{w+s}{s}$  increases and the process is more likely to be chosen by the scheduler.
- Even though short processes are favoured with aging longer process will eventually get through
- Like SRT & SPN need to know expected service time.

# Priority Scheduling

- Priority is associate with each process
- Priorities can be managed with a queue for each priority level
- Dispatcher selects process from the highest priority
- If queue of priority 0.. $m$  are empty and queue  $m+1$  is non-empty, then a process is chosen from queue  $m+1$ .
- This procedure can be used in conjunction with a separate scheduling algorithm for each queue
- Preemptive or non-preemptive?
- SPN is a special case of priority scheduling.



# Priority Boosting

- Major problem in priority scheduling is **starvation**
- A steady stream of higher-priority process can leave some low-priority process waiting indefinitely

RUMOR: When they shut down the IBM7094 at MIT in 1973, they found a low-priority process that had been submitted in 1967 and had not yet been run.

- A solution is **aging**
  - Gradually increase the priority of processes that wait in the system for a long time

# Priority Levels

- Windows 2000 priority levels
  - 32 priority levels from 0 to 31
    - Priority 0 : system level
    - Priority 1 – 15 : variable levels
    - Priority 16 – 31 : real time levels
- Win32 API priority levels
  - 6 priority Class
    - Real-time, high, above normal, normal, below normal and idle
  - 7 relative thread priority levels in each process priority class
    - Time Critical, Highest, Above Normal, Normal, Below Normal, Lowest and Idle




# Priority Boosting

- On completion of I/O operations
- After waiting on executive events or *semaphores*
- After threads in the foreground process complete a wait operation
- When a GUI thread wakes up after windowing activity
- CPU starvation

- SPN, SRT, HRRN can not be used if no indication of the relative length of the processes are available
- We can prefer shorter jobs by penalizing the jobs that have been running for long
- Here we are focusing
  - NOT on the time remaining
  - on the time already executed

# Feedback


- Preemptive scheduling
- For the **FB** policy, we use a set of queues, one for each priority level.



- When a process enters the system
  - it enters the **high priority queue** (*priority  $i$ , with  $i=0$* ).
  - It receives a time slice of service
- When it is pre-empted or blocked,
  - it returns to a lower queue priority queue (*priority  $i+1$* ).
- Within each queue, an FCFS mechanism is used; except for the lowest priority queue, which is round-robin.

# FB Example 1

With equal time slicing:



Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_S$ )	3	6	4	5	2	Mean
<b>FB <math>q = 1</math></b>						
Finish Time	4	20	16	19	11	
Turnaround Time ( $T_T$ )	4	18	12	13	3	10.00
$T_T/T_S$	1.33	3.00	3.00	2.60	1.5	2.29

In this case, behaviour is similar to RR with time quantum of 1.


# FB - Performance

61

- Short process will complete quickly
- Longer process will gradually drift down
- Starvation can occur for long jobs if there are many new jobs (i.e. turnaround time s\_t\_r\_e\_c\_h\_e\_s out).
- **Remedy:** In stead of constant time slices, lower priority queues can get larger time slices (for example, priority  $k$  can get a time slice of  $2^k$  units).

# Feedback Example 2

with time slice of  $2^k$  for queue  $k$ :



Process	A	B	C	D	E	
Arrival Time	0	2	4	6	8	
Service Time ( $T_S$ )	3	6	4	5	2	Mean
<b>FB <math>q = 2i</math></b>						
Finish Time	4	17	18	20	14	
Turnaround Time ( $T_T$ )	4	15	14	14	6	10.60
$T_T/T_S$	1.33	2.50	3.50	2.80	3.00	2.63

# FB - Performance

63

With varying preemption time

- Long processes may still starve.
- **Remedy:** Possible to promote a process to a higher-priority queue after it has waited without service in its current queue.

# Characteristics of scheduling policies

	FCFS	Round robin	SPN	SRT	HRRN	Feedback
<b>Selection function</b>	$\max[w]$	constant	$\min[s]$	$\min[s - e]$	$\max\left\{\frac{w + s}{s}, \frac{w}{s}\right\}$	(see text)
<b>Decision mode</b>	Non-preemptive	Preemptive (at time quantum)	Non-preemptive	Preemptive (at arrival)	Non-preemptive	Preemptive (at time quantum)
<b>Through-Put</b>	Not emphasized	May be low if quantum is too small	High	High	High	Not emphasized
<b>Response time</b>	May be high, especially if there is a large variance in process execution times	Provides good response time for short processes	Provides good response time for short processes	Provides good response time	Provides good response time	Not emphasized
<b>Overhead</b>	Minimum	Minimum	Can be high	Can be high	Can be high	Can be high
<b>Effect on processes</b>	Penalizes short processes; penalizes I/O bound processes	Fair treatment	Penalizes long processes	Penalizes long processes	Good balance	May favor I/O bound processes
<b>Starvation</b>	No	No	Possible	Possible	No	Possible



# Fair-Share Scheduling

- Scheduling algorithms usually individual processes
- Multiple processes might belong to the same application
  - Multiple thread might belong to the same process
- User Point of View:
  - NOT how individual process (thread) is doing
  - How the application (process) is doing
- The idea can even can be extended for a group of independent processes
  - For example processes run by a group



# Summary

- OS must make three types of scheduling decisions
  - Long-term: when new processes are admitted to the system
  - Medium-term: when a process is loaded into main memory
  - Short-term: when a ready process will be executed by the processor
- A variety of criteria is used in designing CPU scheduler
  - User oriented VS System Oriented
  - Quantitative VS Qualitative
- From user point of view response time is usually most important characteristic of a system
- From system point of view throughput or processor utilization is usually more important

# Summary

- Two decision modes for scheduling: Preemptive and Non-preemptive
- Selection may be based on priority, resource requirements, or the execution characteristics of the process
- A variety of algorithms have been developed for short-term scheduling
  - FCFS, RR, SPN, SRT, HRRN, FB, FSS
  - Each has some relative advantage and disadvantage
- The choice of algorithm depends on the expected performance and implementation complexity

# References

- **Operating Systems – Internal and Design Principles**
  - By William Stallings
- Chapter 9