

Application Layer Protocols -2

A/PROF. DUY NGO

Learning Objectives

2.5 P2P applications


2.6 video streaming and content distribution networks

Pure P2P Architecture

- **Not required** a always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:

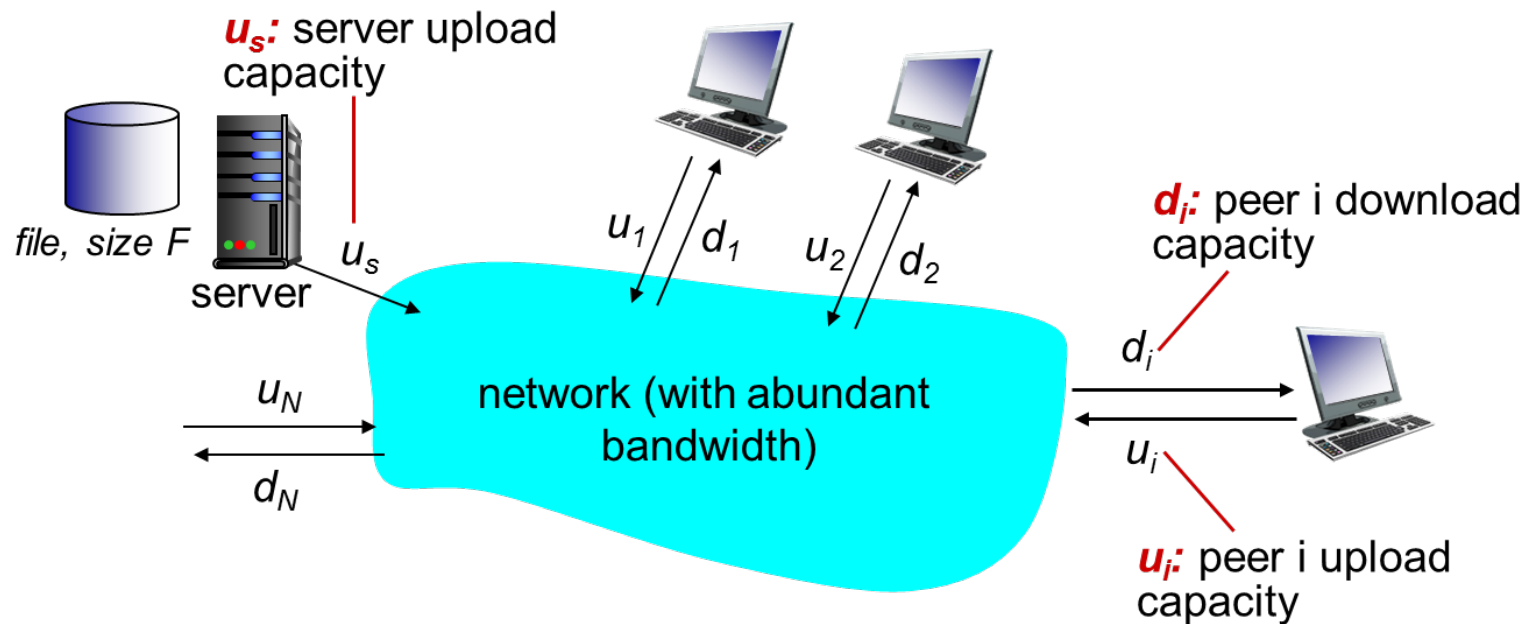
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



File Distribution: Client-Server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?

- peer upload/download capacity is limited resource



File Distribution Time: Client-Server

- **server transmission:** must sequentially send (upload) N file copies:


- time to send one copy: $\frac{F}{u_s}$

- time to send N copies: $\frac{NF}{u_s}$

- **client:** each client must download file copy

- d_{min} = min client download rate

- min client download time: $\frac{F}{d_{min}}$




time to distribute F
to N clients using
client-server approach $D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$

increases linearly in N

File Distribution Time: P2P

- **server transmission:** must upload at least one copy
 - time to send one copy: $\frac{F}{u_s}$
- **client:** each client must download file copy
 - min client download time: $\frac{F}{d_{min}}$
- **clients:** as aggregate must download NF bits



- max upload rate (limiting max download rate) is $u_s + \sum u_i$


time to distribute F
to N clients using
P2P approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{min}, NF/(u_s + \sum u_i)\}$$

increases linearly in N ...
... but so does this, as each peer brings service capacity

Client-Server vs P2P: Example

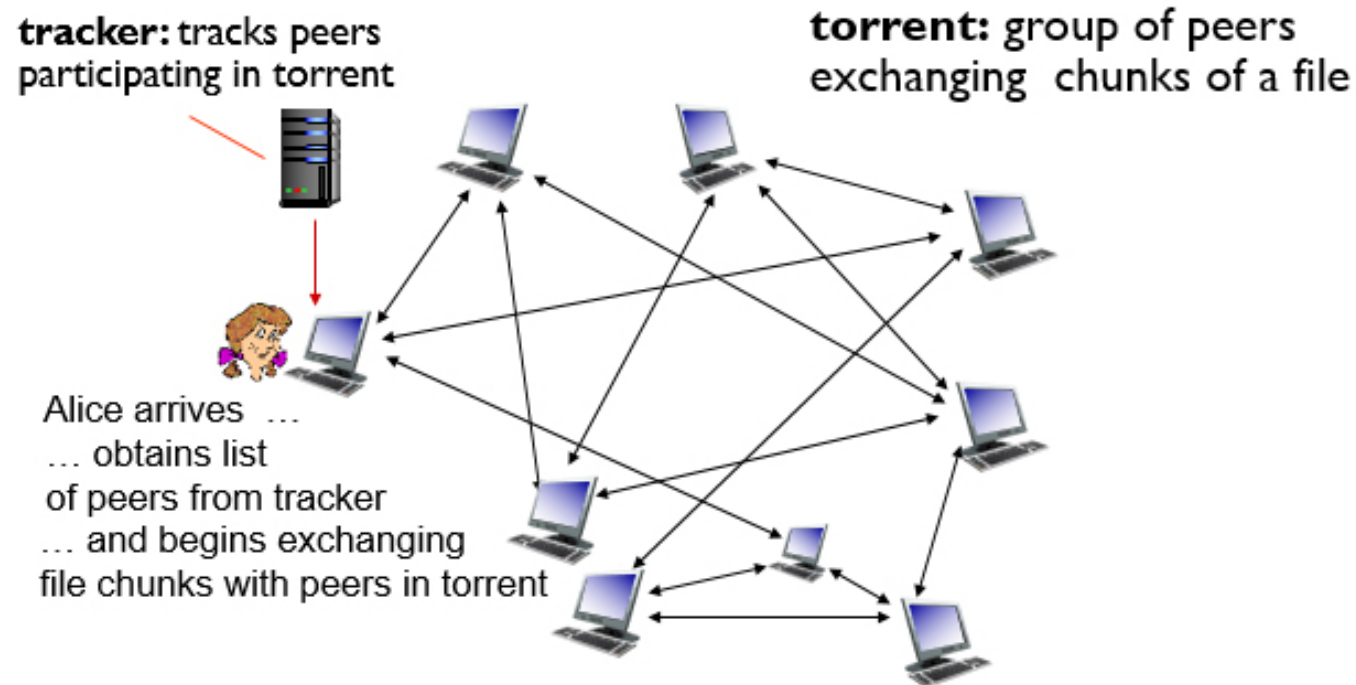
client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$



P2P File Distribution: BitTorrent (1 of 2)


file divided into 256Kb chunks

peers in torrent send/receive file chunks



P2P File Distribution: BitTorrent (2 of 2)

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



BitTorrent: Requesting, Sending File Chunks

requesting chunks:


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first

sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks **at highest rate**
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent: Tit-For-Tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Video Streaming and CDNs: Context

- video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **solution:** distributed, application-level infrastructure



Multimedia: Video (1 of 2)

video: sequence of images displayed at constant rate

- e.g., 24 images/sec

digital image: array of pixels

- each pixel represented by bits

coding: use redundancy **within** and **between** images to decrease # bits used to encode image

- spatial (within image)
- temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

Multimedia: Video (2 of 2)

CBR: (constant bit rate): video
encoding rate fixed

VBR: (variable bit rate): video
encoding rate changes as amount of
spatial, temporal coding changes

examples:

- MPEG1 (CD-ROM) 1.5 Mbps
- MPEG2 (DVD) 3-6 Mbps
- MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i


temporal coding example:
instead of sending complete frame at $i+1$, send only differences from frame i



frame i+1

Streaming Stored Video

simple scenario:



Streaming Multimedia: DASH (1 of 2)

DASH: **D**ynamic, **A**daptive **S**treaming over **H**TT**P**

server:

- divides video file into multiple chunks
- each chunk stored, encoded at different rates
- **manifest file:** provides URLs for different chunks

client:

- periodically measures server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

Streaming Multimedia: DASH (2 of 2)

“intelligence” at client: client determines

- **when** to request chunk (so that buffer starvation, or overflow does not occur)
- **what encoding rate** to request (higher quality when more bandwidth available)
- **where** to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

Content Distribution Networks (1 of 2)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of **simultaneous** users?

option 1: single, large “mega-server”

- single point of failure
- point of network congestion
- long path to distant clients
- multiple copies of video sent over outgoing link

....quite simply: this solution **doesn't scale**

Content Distribution Networks (2 of 2)

option 2: store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)

- **enter deep:** push CDN servers deep into many access networks
 - close to users
 - used by Akamai, 1700 locations
- **bring home:** smaller number (10's) of larger clusters in POPs near (but not within) access networks
 - used by Limelight


Content Distribution Networks (CDNs) (1 of 2)

CDN: stores copies of content at CDN nodes

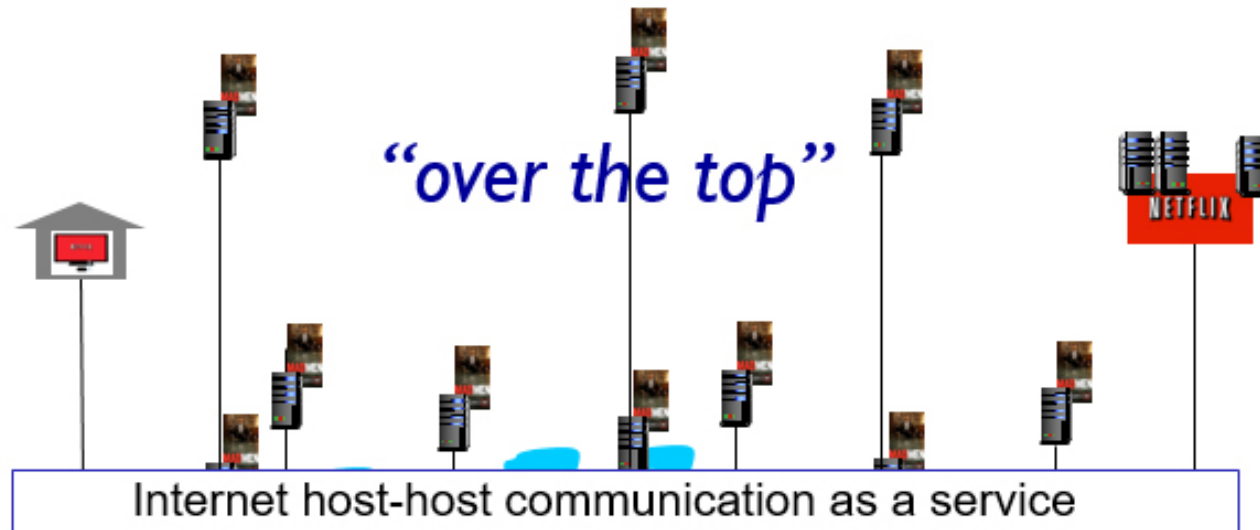
- e.g. Netflix stores copies of MadMen

subscriber requests content from CDN

- directed to nearby copy, retrieves content
- may choose different copy if network path congested



Content Distribution Networks (CDNs) (2 of 2)



OTT challenges: coping with a congested Internet


- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

CDN Content Access: A Closer Look

Bob (client) requests video <http://netcinema.com/6Y7B23V>

– video stored in CDN at <http://KingCDN.com/NetC6y&B23V>

1. Bob gets URL for video <http://netcinema.com/6Y7B23V> from netcinema.com web page
2. resolve <http://netcinema.com/6Y7B23V> via Bob's local DNS
3. netcinema's DNS returns URL <http://KingCDN.com/NetC6y&B23V>
- 4&5. Resolve <http://KingCDN.com/NetC6y&B23V> via KingCDN's authoritative DNS, which returns IP address of KingCDN server with video



Case Study: Netflix

