

# Transport Layer Protocols -4

---

A/PROF. DUY NGO

# Learning Objectives

---

**3.6** principles of congestion control

**3.7** TCP congestion control

# Principles of Congestion Control

---

## **congestion:**

informally: “too many sources sending too much data too fast for **network** to handle”

different from flow control!

manifestations:

- lost packets (buffer overflow at routers)
- long delays (queueing in router buffers)

a top-10 problem!

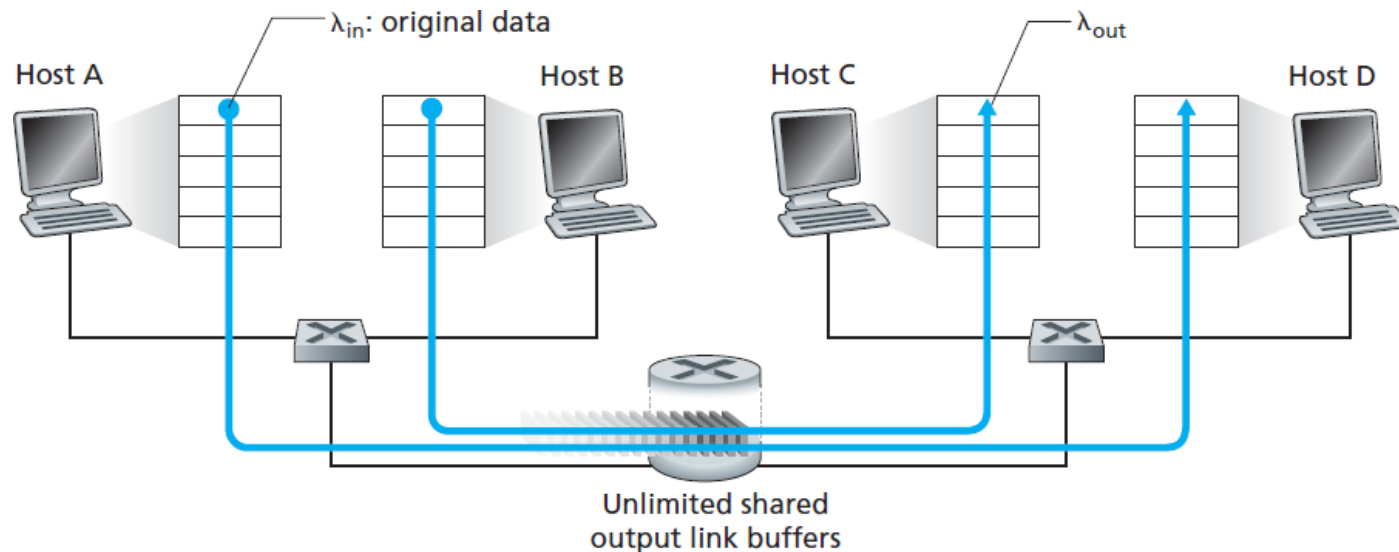
# Causes/Costs of Congestion: Scenario 1 (1 of 2)

two senders, two receivers

one router, infinite buffers


output link capacity:  $R$

no retransmission




# Causes/Costs of Congestion: Scenario 1 (2 of 2)

- maximum per-connection throughput:  $\frac{R}{2}$




- large delays as arrival rate,  $\lambda_{in}$ , approaches capacity



# Causes/Costs of Congestion: Scenario 2 (1 of 6)

- one router, **finite** buffers
- sender retransmission of timed-out packet
  - application-layer input = application-layer output:
  - transport-layer input includes **retransmissions** :


$$\lambda_{in} = \lambda_{out} \quad \lambda'_{in} \geq \lambda_{in}$$



# Causes/Costs of Congestion: Scenario 2 (2 of 6)

## idealization: perfect knowledge


- sender sends only when router buffers available



# Causes/Costs of Congestion: Scenario 2 (3 of 6)

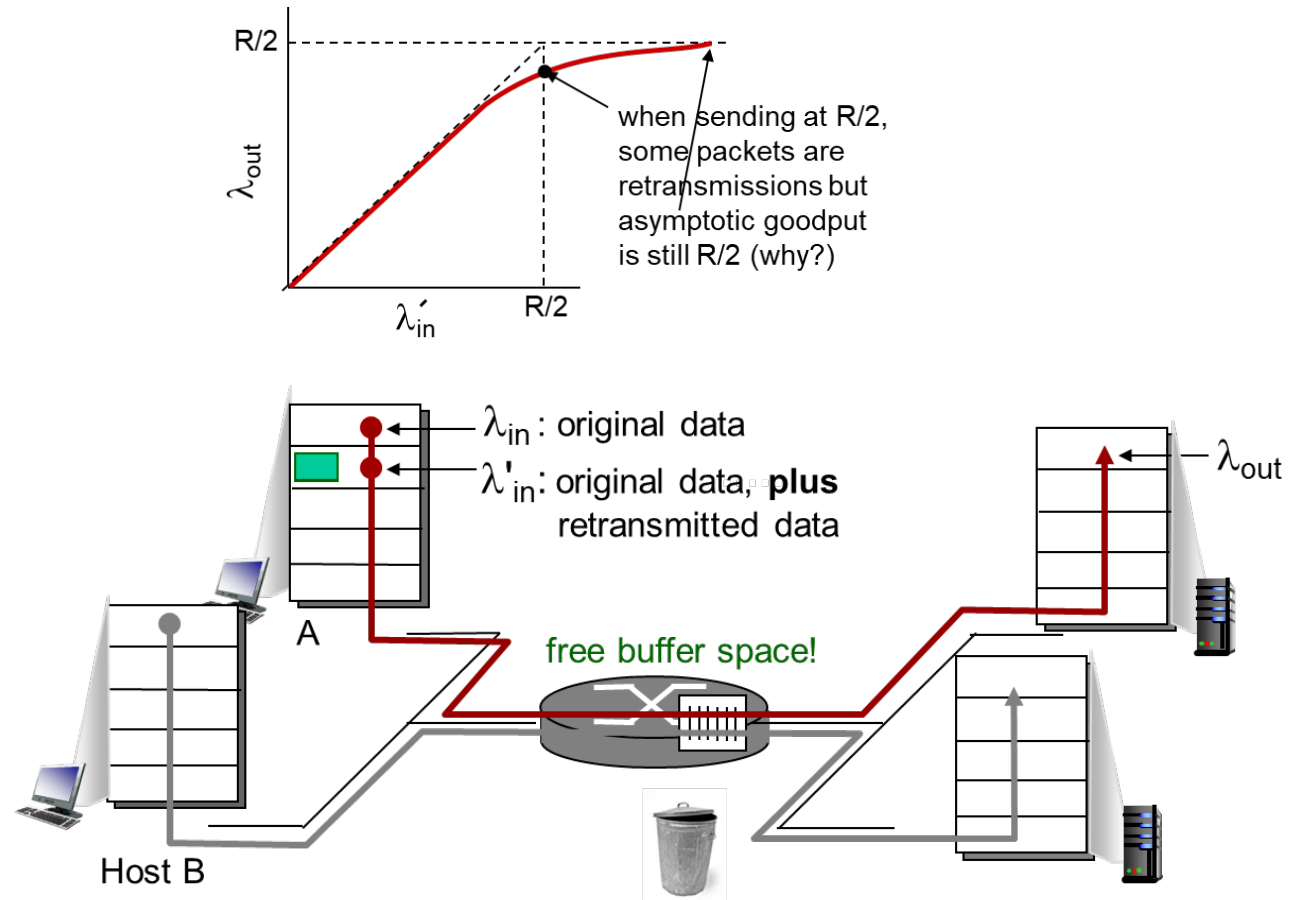
**Idealization: known loss** packets can be lost, dropped at router due to full buffers

sender only resends if packet **known** to be lost





# Causes/Costs of Congestion: Scenario 2 (4 of 6)




# Causes/Costs of Congestion: Scenario 2 (5 of 6)

## Realistic: duplicates

packets can be lost, dropped at router due to full buffers

sender times out prematurely, sending **two** copies, both of which are delivered




# Causes/Costs of Congestion: Scenario 2 (6 of 6)

## “costs” of congestion:

more work (retrans) for given “goodput”

unnneeded retransmissions: link carries multiple copies of pkt

- decreasing goodput




# Causes/Costs of Congestion: Scenario 3 (1 of 2)



- four senders
- multihop paths
- timeout/retransmit

**Q:** what happens as  $\lambda_{in}$  and  $\lambda'_{in}$  increase ?

**A:** as red  $\lambda'_{in}$  increases, all arriving blue pkts at upper queue are dropped, blue throughput  $\rightarrow 0$



# Causes/Costs of Congestion: Scenario 3 (2 of 2)



## another “cost” of congestion:


when packet dropped, any “upstream transmission capacity used for that packet was wasted!

# TCP Congestion Control: Additive Increase Multiplicative Decrease


**approach:** sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs

- **additive increase:** increase `cwnd` by 1 MSS every RTT until loss detected
- **multiplicative decrease:** cut `cwnd` in half after loss

AIMD saw tooth  
behavior: probing  
for bandwidth



# TCP Congestion Control: Details



- sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- **cwnd** is dynamic, function of perceived network congestion


## TCP sending rate:

- **roughly:** send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

# TCP Slow Start

- when connection begins, increase rate exponentially until first loss event:
  - initially `cwnd` = 1 MSS
  - double `cwnd` every RTT
  - done by incrementing `cwnd` for every ACK received
- **summary:** initial rate is slow but ramps up exponentially fast





# TCP: Detecting, Reacting to Loss

---

loss indicated by **timeout**: **TCP RENO**

- **cwnd** set to 1 MSS;
- window then grows exponentially (as in Slow Start) to threshold, then grows linearly

loss indicated by **3 duplicate ACKs**: **TCP RENO**

- dup ACKs indicate network capable of delivering some segments
- **cwnd** is cut in half window plus 3 (for Fast Recovery), then grows linearly

**TCP Tahoe** always sets **cwnd** to 1 (for both timeout and 3 duplicate acks) and enter Slow Start


# TCP: Switching from Slow Start to CA

**Q:** when should the exponential increase switch to linear?

**A:** when **cwnd** gets to  $\frac{1}{2}$  of its value before timeout.

**Implementation:**

- variable **ssthresh**
- on loss event, **ssthresh** is set to  $\frac{1}{2}$  of **cwnd** just before loss event



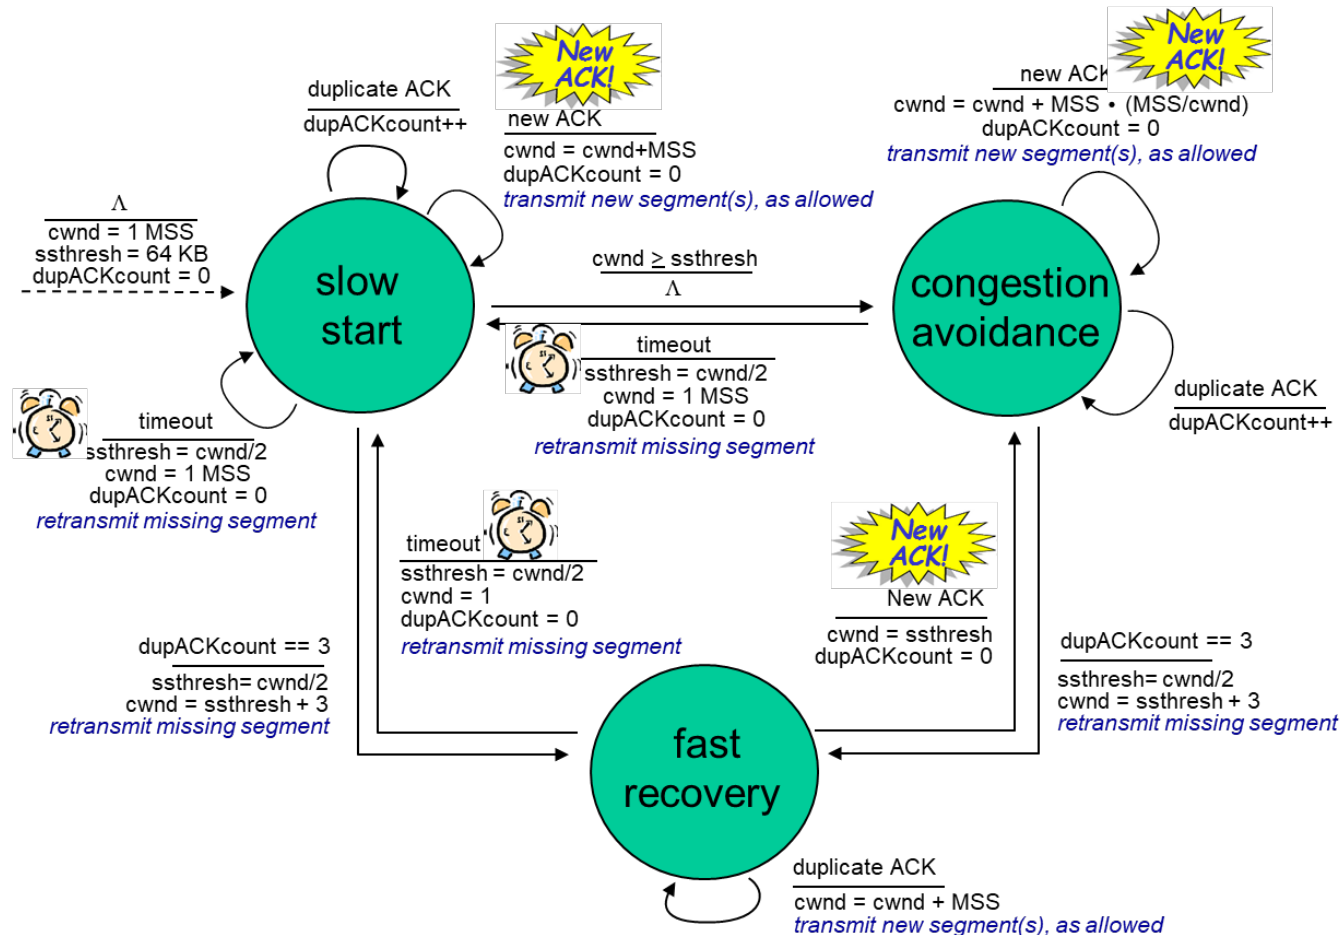
**Figure 3.52** ♦ Evolution of TCP's congestion window (Tahoe and Reno)

In Fig 3.52, at transmission round 9, triple duplicate ACKs occur

- $\text{ssthresh} := \text{cwnd}/2 = 12/2 = 6 \text{ MSS}$
- **Reno:**  $\text{cwnd} := \text{ssthresh} + 3 = 6 + 3 = 9 \text{ MSS}$
- **Tahoe:**  $\text{cwnd} := 1 \text{ MSS}$

\* Check out the online interactive exercises for more examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)


# Summary: TCP Congestion Control



# TCP Throughput

- average TCP thruput as function of window size, RTT?
  - ignore slow start, assume always data to send
- W: window size (measured in bytes) where loss occurs
  - average window size (# in-flight bytes) is
  - average thruput is  $\frac{3}{4}W$  per RTT  $\frac{3}{4}W$

$$\text{avg TCP thruput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes / sec}$$



# TCP Futures: TCP over “Long, Fat Pipes”

---

- example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- requires  $W = 83,333$  in-flight segments
- throughput in terms of segment loss probability,  $L$  [Mathis 1997]:

$$\text{TCP throughput} = \frac{1.22 \cdot \text{MSS}}{\text{RTT} \sqrt{L}}$$


→ to achieve 10 Gbps throughput, need a loss rate of

$$L = 2 \cdot 10^{-10} \quad \text{– a very small loss rate!}$$

- new versions of TCP for high-speed

# TCP Fairness

**fairness goal:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have average rate of  $\frac{R}{K}$




# Why is TCP Fair?

---

two competing sessions:

additive increase gives slope of 1, as throughput increases

multiplicative decrease decreases throughput proportionally



# Fairness (More)

---

## Fairness and UDP

- multimedia apps often do not use TCP
  - do not want rate throttled by congestion control
- instead use UDP:
  - send audio/video at constant rate, tolerate packet loss

## Fairness, parallel TCP connections

- application can open multiple parallel connections between two hosts
- web browsers do this
- example, link of rate  $R$  with 9 existing connections:
  - new app asks for 1 TCP, gets rate  $\frac{R}{10}$
  - new app asks for 11 TCPs, gets  $\frac{R}{2}$



# Explicit Congestion Notification (ECN)

## network-assisted congestion control:

- two bits in IP header (ToS field) marked **by network router** to indicate congestion
- congestion indication carried to receiving host
- receiver (seeing congestion indication in IP datagram) ) sets ECE bit on receiver-to-sender ACK segment to notify sender of congestion

