

Priority queues and sorting

Priority queues

Compositions

Insertion sort

Selection sort

Heap sort

Read Chapter 18 of the textbook!



Priority queues

- A priority queue is an ADT for storing a collection of elements
- The elements may be arbitrarily inserted but removal is in order of priority
 - priority is specified using a parameter which is defined for each element stored in the collection
 - the parameter used to implement priority is called the **priority** or **key** for the element
 - priorities are not unique, and may be changed if necessary
 - elements with the highest priority will be removed first
 - there may be multiple elements with the same priority
 - in which case it is common to form a FIFO queue of those elements

Priority queues



- Example: **Waiting queue for an event**
 - Suppose the organisers of an event will sell tickets following a priority queue, aiming at VIP attendees.
 - The event is sponsored by company X, which has a loyalty program.
 - Clients who call to express interest for the event are put in a waiting queue. When sales start, they will be called by order of priority.
 - The priority follows the number of loyalty points accumulated (not the time that they called and were put in the queue).
 - If two customers have the same number of loyalty points, the tie is broken by the time of the call.
-
- How would you implement that?
 - Which data structure would you use?
 - What would be the object stored composed of?
 - Think about it for 5 minutes.

Priority queues

- A priority queue is a container for which every element has a priority, assigned at the time of insertion, that provides a means of selecting the next element to be removed
- A priority queue `PQueue` extends the `Queue` class by adding the following fundamental method:

```
// insert element elem with priority k into PQueue
void enqueue(Key k, value_type elem)
```
- Of course, `dequeue` will also change, so that when it is called, the element with the highest priority is removed
- Note that only the element should be returned, not its priority (unless that is an explicit requirement)

Using priority queues for sorting



- A set of n elements can be sorted by inserting them into a priority queue P and then calling `dequeue()` until P is empty
- The algorithm is:
 - put the n elements of the set into an empty priority queue using n enqueue operations
 - extract the elements from P using n dequeue operations
- For example, if the elements were stored in a sequence, they could be removed from the sequence into the priority queue, then added back to the sequence, thus sorting the elements in the sequence
- This is the underlying method used by popular sorting algorithms such as selection sort, insertion sort and heap sort

Priority queue ADT

- Using a linked list as the underlying data structure

```
// Adds element to the tail of the queue with a default priority
void enqueue(value_type elem)
```

```
// Adds element to the tail of the queue with specified priority
void enqueue(int priority, value_type elem)
```

```
// Removes the element with the highest priority
value_type dequeue()
```

- Now, how do we keep the element and its priority connected?
- Use a composition pattern

Priority queue ADT

- The composition pattern defines a single object c , which is the composition of other objects
- The objects stored in the queue will be pairs (key, element) implemented as a new class

```
template <typename value_type>
class Comp
{
public:
    Comp(Key k, value_type e) {itemKey = k; elem = e;} // Constructor

    // Query member functions
    Key get_key() {return itemKey;}
    value_type get_element() {return elem;}

    // Mutator member functions
    void set_key(Key k) {itemKey = k;}
    void set_element(value_type e) {elem = e;}

private:
    Key itemKey;
    value_type elem;
};
```

Selection sort



- Consider a `PriorityQueueSort` method with the following declaration:

```
// Receives an array and orders it by priority
void PriorityQueueSort(value_type input[])
```

- The method:
 - takes an unsorted array of n elements
 - inserts the elements into a priority queue
 - repeatedly removes the minimum element from the priority queue, adding them to the original array, until the priority queue is empty
- The ordering of elements is done at the time of removal from the priority queue (i.e. when they are selected to be removed).
- Insertion of elements takes constant $O(1)$ time (e.g. at the tail of the linked list). Selection of elements take $O(n)$, i.e. the whole linked list has to be traversed to determine the one with the highest priority.

Insertion sort



- Consider a PriorityQueueSort method with the following declaration:

```
// Receives an array and orders it by priority  
void PriorityQueueSort(value_type input[])
```
- The method:
 - takes an unsorted array of n elements
 - inserts the elements into a priority queue
 - repeatedly removes the minimum element from the priority queue, adding them to the original array, until the priority queue is empty
- The ordering of the elements is done at the time of insertion into the priority queue.
- Insertion of elements takes $O(n)$ time, i.e. the linked list has to be traversed to determine the correct point of insertion to keep it ordered. Removal takes constant $O(1)$ time (e.g. remove from the head of the linked list).

Complexity



- If you are trying to order n items using insertion sort, each insertion will take time $O(n)$. Since you have n elements in total, the complexity then becomes $O(n*n) = O(n^2)$
- If you are trying to order n items using selection sort, each removal will take time $O(n)$. Since you have n elements in total, the complexity is also $O(n^2)$
- This complexity is relatively high. A complexity of $O(n^2)$ indicates that the number of operations (comparisons, assignments, etc) required grows with the square of the size of the array.
- E.g. if sorting an array takes 10 seconds, sorting an array $2x$ as long will take 40 seconds. This level of computational complexity makes insertion sort and selection sort unsuitable for high performance applications.

Heap sort

- Heap sorting has a better performance compared to selection and insertion sorting, as instead of a linear data structure (linked list), it uses a binary tree.
- The heap used is a binary tree that:
 - stores item-key pairs at its internal nodes and null at its leaf nodes
 - satisfies the following properties:
 - for every internal node other than the root, the key stored at the node is less than or equal to the key stored at the node's parent
 - thus the element with the largest key will always be stored at the root node
 - is complete
 - all levels, except for the lowest level, have the maximum number of nodes
 - the lowest level is filled from the left to the right

Implementing sorting using a heap



- The priority queue P has instance data comprising:
 - a heap T for which each internal node v stores
 - an element of the priority queue
 - the key k for that element
 - a reference to the position of the last node of T
- The heap is itself implemented using one of the binary tree data structures presented previously
- The elements stored in the binary tree are key-item pairs

Enqueuing using a heap



- To implement the method `enqueue` of the priority queue ADT using T
 - a new internal node is added to T
 - this becomes the new last node of T
 - the new internal node is created by selecting the appropriate external node z and calling `expandExternal(z)` on it
 - This makes z an internal node with two `NULL` child nodes
 - z is usually the external node immediately to the right of the last insertion position
 - Exceptions to this occur when either
 - the current last node is the right-most node on its level, in which case z is the left-most node of the bottom level, or
 - there are no internal nodes in which case z is the root node of T

Enqueuing using a heap



- after `expandExternal(z)` finishes
 - z becomes the new last node, and
 - the new data element-key pair (k, e) is stored there, i.e. $k(z) = k$
- at this stage T is complete, but does not necessarily comply with the heap-order property
- unless z is the root of T , we need to compare $k(z)$ with the key of z 's parent u
 - if $k(u) < k(z)$ (in the case of "maximum root node") then heap-order has been violated by the insertion and the key-element pairs stored at u and z need to be swapped
 - this may violate heap-order higher up the tree, so swapping is continued until heap-order is not violated

Enqueuing using a heap



- The process of moving the new key-element pair into its correct position in the tree is called *up-heap bubbling*
 - in the worst case the bubbling causes the new key-element pair to move up to the root position
 - the cost of up-heap bubbling is proportional to the height of the tree, and since the tree height is given by $h = \log_2(n + 1)$ the cost of up-heap bubbling is $O(\log n)$
 - E.g. for an array of 1,023 elements, the associated heap will be a binary tree with 10 levels
- To start the process of heap insertion it is necessary to find the insertion position z
 - the heap keeps as instance data the last insertion position w

Enqueuing using a heap



- To find the insertion position using the binary tree ADT:
 - start at w and move up the tree calling the `parent()` function until either the root (determined by using `isRoot()`) or w is found to be a left-child (determined by using `isLeftChild`)
- The method `isLeftChild` would be implemented as follows:

```
bool isLeftChild(Position p)
{
    if (isRoot(p)) {return false;}
    else
        return (leftChild(parent(p)) == p);
}
```

Enqueuing using a heap



- If the root has been reached, then the last insertion position was the right-most internal node of its level
 - set u to the root
- If the parent of a left-child was reached
 - set u to be the sibling of the left-child (i.e. the right-child of the reached node's parent)
- Starting at u , move down the tree using `leftChild()` until an external node z is reached
 - this may involve zero or more calls to `leftChild()`
 - Insertion code follows on the next slide.

Enqueuing using a heap

- Insertion code

```
void insertItem(Key k, Item e)
{
    Position z; // Position to insert
    if (isEmpty())
        z = root();
    else
    {
        z = last;
        while (!isRoot(z) && !isLeftChild(z))
            z = parent(z);
        if (!isRoot(z))
            z = rightChild(parent(z));
        while (!isExternal(z))
            z = leftChild(z);
    }
    expandExternal(z);
    replace(z, new Comp(k, e));
    last = z;
    ...
}
```

Enqueuing using a heap

- Insertion code (up heap bubbling)

```
// Up-heap bubbling comes next (max at root node)
...
Position u;

while (!isRoot(z)) //Up-heap bubble (z is the last
{ //insertion position
    u = parent(z);
    if (z.get_comp().get_key() < u.get_comp().get_key()) {break;}
    swap(u, z);
    z = u;
}
```

Dequeuing using a heap



- The procedure for locating the new last node in an insertion can be reversed to update the last node after removal
- Removal from the heap is necessary to implement the dequeue method of the priority queue class
- It is known that the maximum element is stored at the root of the heap T
- unless the root is the only internal node we cannot simply delete the root node, because that would ruin the binary tree
- what we do is copy the key-element pair stored at the last node w into the root node
- then we delete the last node using `removeExternal(w)`

Dequeuing using a heap



- The procedure for locating the new last node in an insertion can be reversed to update the last node after removal
- Removal from the heap is necessary to implement the dequeue method of the priority queue class
- It is known that the maximum element is stored at the root of the heap T
- unless the root is the only internal node we cannot simply delete the root node, because that would ruin the binary tree
- what we do is copy the key-element pair stored at the last node w into the root node
- then we delete the last node using `removeExternal(w)`

Example



- Sorting a small vector

[5 1 8 4 10 3 6 9]

- Add all elements to a heap, and then remove them sequentially.

See you next week!

