



COMP1140: Database and Information Management

Lecture Note – Week 11

Dr Suhuai Luo

School of SEEC

University of Newcastle



Note

- Pls take your time to fill SFC, Thanks. You can do it
 - Either follow the link from your email notification, or
 - Goto myuon.newcastle.edu.au , or
 - the homepage of UONLine (Blackboard)
- SQL test result will be released on BB from Tue 12pm. You can discuss with your instructor on your performance in this week lab only.
- Assignment 3 is due next week at the beginning of your lab session
 - Deliver hard copy to your marker; softcopy to BB as well
 - You are required to attend your lab session, the SQL part of the A3 will be marked on site – no show, no mark
 - Read the assignment specifications & the marking guide; make sure to implement normalised tables; meet constraints etc.
 - More Q & discussion at the end of this lecture
- Course review will be done in lecture next week. Pls come.



Note

- Pls take your time to fill SFC, Thanks. You can do it
 - Either follow the link from your email notification, or
 - Goto myuon.newcastle.edu.au , or
 - the homepage of UONLine (Blackboard)
- SQL test result will be released on BB from Tue 12pm. You can discuss with your instructor on details of your performance, in your lab this week only.
- Assignment 3 is due next week at the beginning of your lab session
 - Deliver hard copy to your marker; softcopy to BB as well
 - You are required to attend your lab session, the SQL part of the A3 will be marked on site – no show, no mark
 - Read the assignment specifications & the marking guide; make sure to implement normalised tables; meet constraints etc.
 - More Q & discussion at the end of this lecture
- Course review will be done in lecture next week. Pls come.



Last Lecture

- Views
- Transactions
- Triggers
- Stored Procedures



This week

- Discretionary Access Control

- Schema
- Privileges
- Users and Roles
- GRANT, DENY and REVOKE

Ref. chapter 20.2, chapter 7, chapter 8

- Physical Database Design

- Data Storage: Overview
- Indexes: Overview
- Query Execution Plans: Overview

Ref. chapter 18, appendix F

Database authorization mechanisms



- Mandatory access control - is based on system-wide policies that cannot be changed by individual users
 - Each database object is assigned a certain classification level
 - Each subject is given a designed clearance level
 - A subject requires necessary clearance to read/write a DB object
- Discretionary Access Control - each user is given appropriate access rights/privileges on specific database objects
 - Typically users obtain certain privileges when they create an object and can pass some or all of these privileges to other users
 - SQL supports discretionary access control through statements GRANT & REVOKE



SQL: Naming database objects

- Database objects are referenced via a hierarchy:
 - Server-Database-Schema-Object
- Each database object has a **fully qualified name** as follows:
 - *ServerName.DatabaseName.SchemaName.ObjectName*
 - The fully qualified name is unique for each database object
 - **ServerName**: Name of the database server
 - **Database**: Name of database within the server
 - **SchemaName**: Name of schema within the database
 - **ObjectName**: Name of the database object



Database Schema

- Relations and other database objects exist in an *environment*.
- Each environment contains one or more *catalogs*, and each catalog consists of set of schemas.
- Schema is named collection of related database objects.
- Objects in a schema can be tables, views, domains, collations, translations, and character sets. All have same owner



CREATE SCHEMA

- CREATE SCHEMA [Name | AUTHORIZATION CreatorId]
- DROP SCHEMA Name [RESTRICT | CASCADE]
- With RESTRICT (default), schema must be empty or operation fails.
- With CASCADE, operation cascades to drop all objects associated with schema in the order defined when created. If any of these operations fail, DROP SCHEMA fails.



CREATE SCHEMA (cont'd)

Example:

```
CREATE SCHEMA Sprockets  
AUTHORIZATION Annik
```

```
CREATE TABLE NineProngs (source int,  
cost int, partnumber int)
```

```
GRANT SELECT TO Mandar
```

```
DENY SELECT TO Prasanna;
```

```
GO
```



Access Control - Authorization Identifiers and Ownership

- Authorization identifier is normal SQL identifier used to establish identity of a user.
- Usually has an associated password.
- Used to determine which objects user may reference and what operations may be performed on those objects.
- Each object created in SQL has an owner, as defined in `AUTHORIZATION` clause of schema to which object belongs.
- Owner is the only person who may know about the object.



Privileges

- Are the actions that a user is permitted to carry out on a given base table or view
- Common DML privileges
 - **SELECT**: Retrieve data from table/view.
 - **INSERT**: Insert new rows into table/view.
 - **UPDATE**: Modify rows of data in a table/view.
 - **DELETE**: Delete rows of data from a table/view.
 - **REFERENCES**: Reference columns of named table in integrity constraints.
 - **USAGE**: use domains, collations, character sets



Privileges (contd.)

- Common DDL privileges
 - **ALTER**: Confers the ability to change the properties of object (ability to alter, create, or drop any securable that is contained within that scope).
 - For example, ALTER permission on a schema includes the ability to create, alter, and drop objects from the schema
 - **ALTER ANY**: Confers the ability to CREATE, ALTER, or DROP individual instances of the objects.
 - For example, ALTER ANY SCHEMA confers the ability to create, alter, or drop any schema in the database



Privileges (contd.)

- Privileges may be granted at different levels
 - Server level
 - Database level
 - Schema level
 - Object level



Privileges (contd.)

- Examples of database-level privileges:
 - **BACKUP DATABASE**: Ability to backup a database
- Example of server-level privileges
 - **ALTER ANY DATABASE**: Ability to create, drop or configure any database



Granting Privileges

- SQL has the GRANT statement used for granting permissions:
- Format:
GRANT {PrivilegeList | ALL PRIVILEGES}
ON ObjectName
TO {AuthorizationIdList | PUBLIC}
[WITH GRANT OPTION]



Granting Privileges (contd.)

- *PrivilegeList* consists of one or more of above privileges separated by commas.
- ALL PRIVILEGES grants all privileges to a user.
- PUBLIC allows access to be granted to all present and future authorized users.
- *ObjectName* can be a base table, view, domain, character set, collation or translation.
- WITH GRANT OPTION allows privileges to be passed on.



Example: GRANT

- Give John White full privileges to Staff table.

GRANT ALL PRIVILEGES

ON Staff

TO JohnWhite WITH GRANT OPTION;

- Give users Personnel and Director SELECT and UPDATE on column salary of Staff.

GRANT SELECT, UPDATE (salary)

ON Staff

TO Personnel, Director;



Example: GRANT Specific Privileges to PUBLIC

- Give all users SELECT on Branch table.

```
GRANT SELECT  
ON Branch  
TO PUBLIC;
```



Managing Privileges with Roles

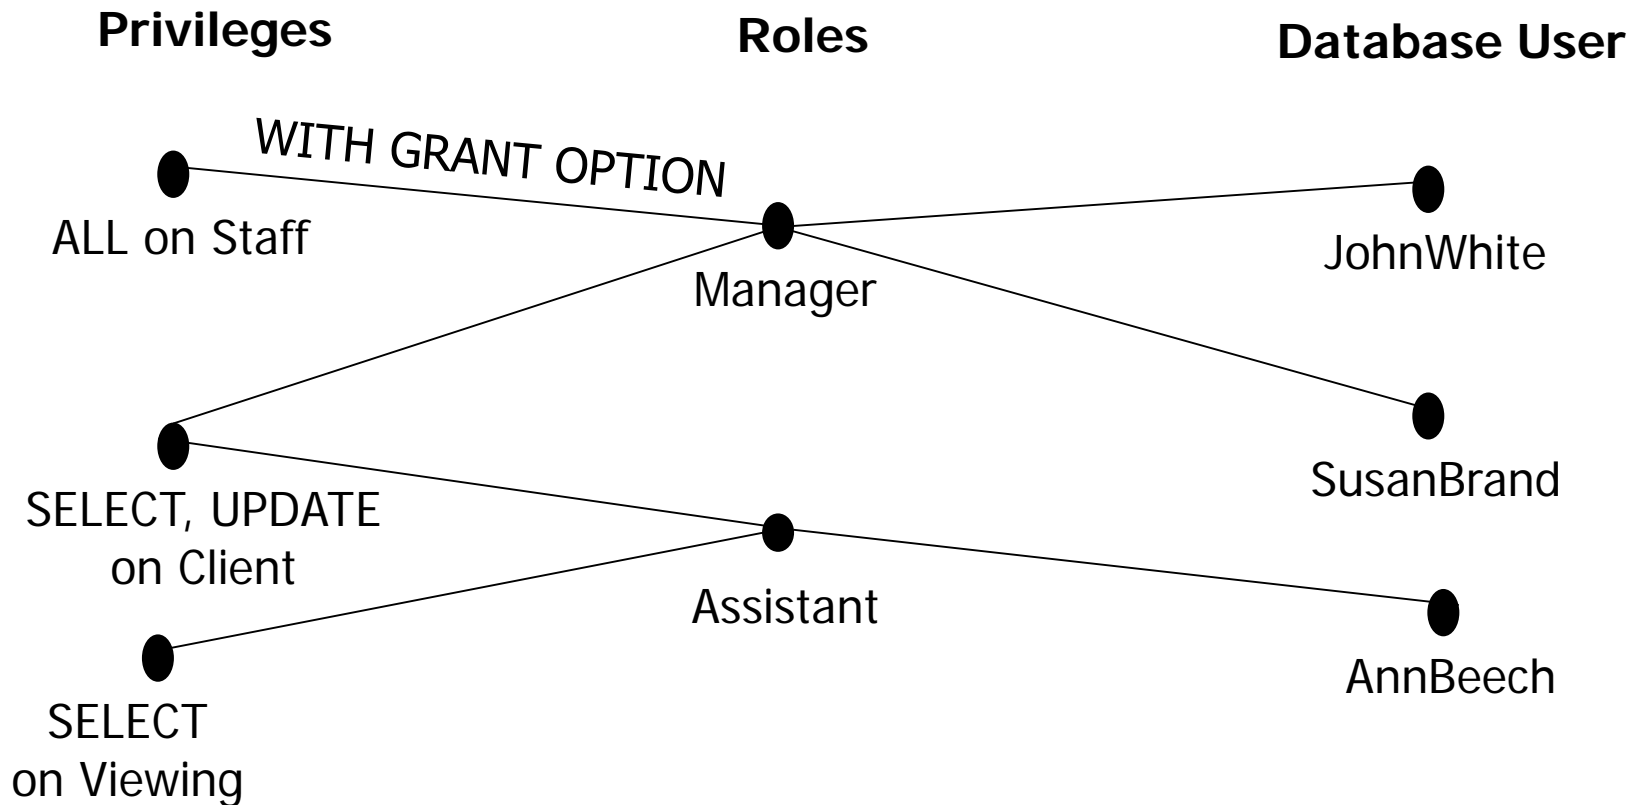
- It becomes quite arduous to manage a large number of users and permissions.
- *Roles* enable group permission and manage it easier.
- Example,
 - CREATE ROLE Manager;



Roles (contd.)

- Now permissions allowed to *managers* can be provided to the Manager role.
 - GRANT ALL PRIVILEGES
ON Staff
TO Manager
- By adding database user “John White” as a member of the Manager role, JohnWhite will acquire all privileges of *Manager* role

Roles (contd.)





Roles (contd.)

- Managing permission would be convenient
 - Add a permission to all managers => Add a permission to role manager
 - A new manager is recruited. => Add new database user to role managerEtc.
- Drop a role with DROP ROLE statement



DENY

- Denies a permission to a principal.
- Prevents that principal from inheriting the privilege through its role memberships
- Basic Syntax:
 - DENY { ALL [PRIVILEGES] | *permission* }
 - ON ObjectName
 - TO AuthorizationIdList



Example: DENY

- Deny SELECT privilege on Viewings table to JohnWhite (even inheriting the privilege through role membership)

```
DENY SELECT  
ON Viewings  
TO JohnWhite
```



REVOKE

- **REVOKE** removes a previously granted or denied permission

REVOKE [GRANT OPTION FOR]

{PrivilegeList | ALL PRIVILEGES}

ON ObjectName

FROM {AuthorizationIdList | PUBLIC}

[RESTRICT | CASCADE]

- ALL PRIVILEGES refers to all privileges granted to a user by the user who is revoking the privileges



Example: REVOKE Specific Privileges

Revoke privilege SELECT on Branch table from all users.

```
REVOKE SELECT  
ON Branch  
FROM PUBLIC;
```

Revoke all privileges given to Director on Staff table.

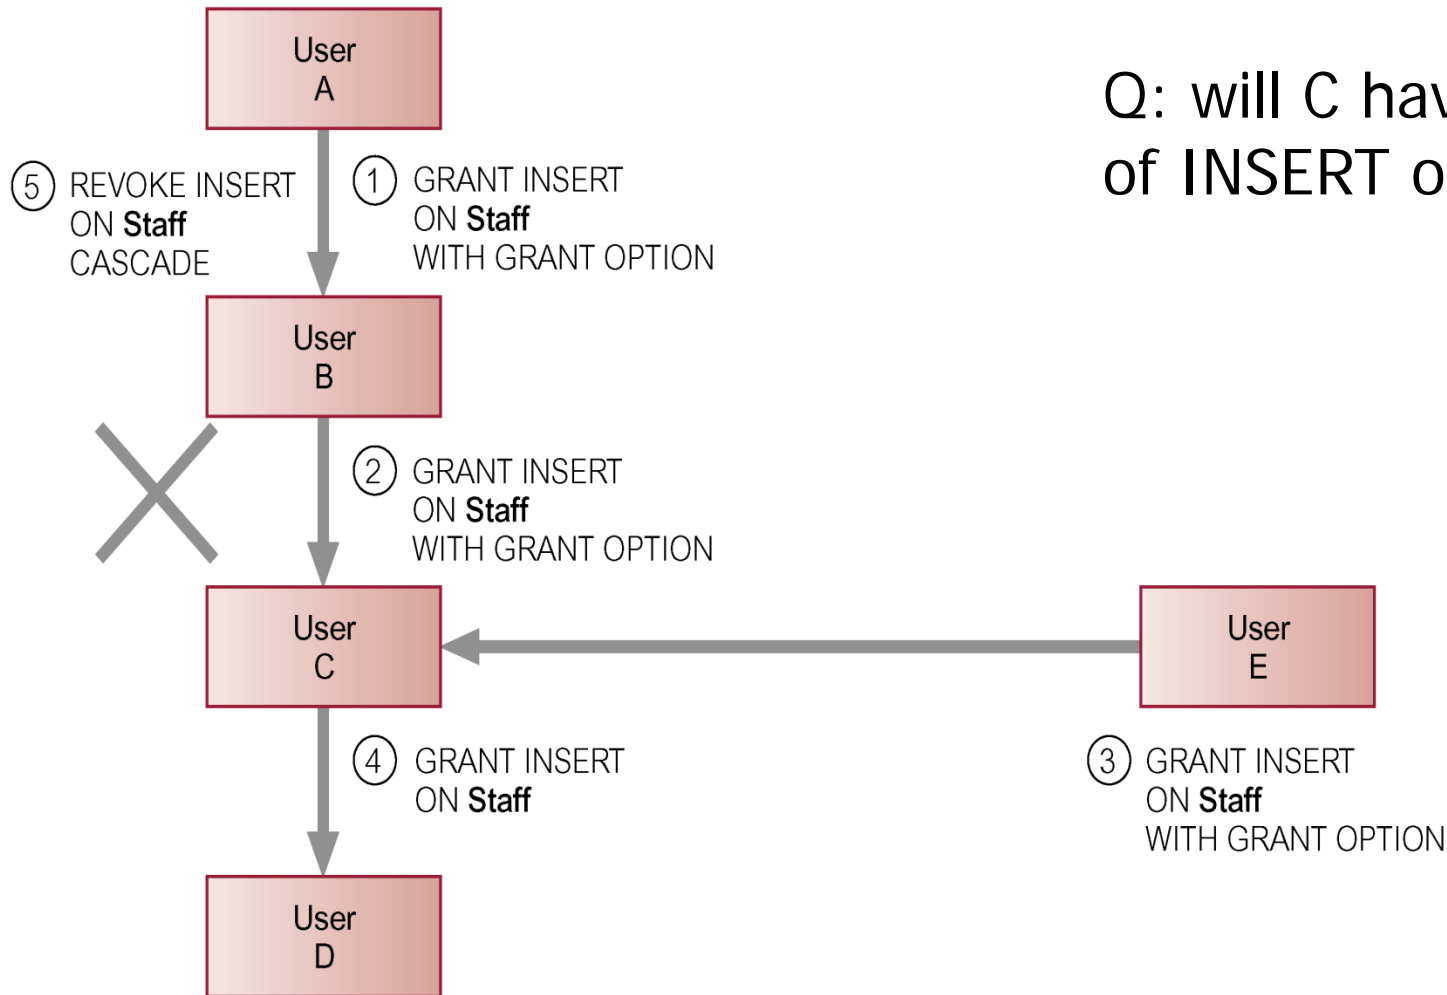
```
REVOKE ALL PRIVILEGES  
ON Staff  
FROM Director;
```



REVOKE (contd.)

- GRANT OPTION FOR allows privileges passed on via WITH GRANT OPTION of GRANT(i.e., the right to grant to 3rd party) to be revoked separately from the privileges themselves.
- REVOKE fails if it results in an abandoned object, such as an abandoned view, unless the CASCADE keyword has been specified.
- Privileges granted to this user by other users are not affected.

REVOKE (contd.)





Summary

- Database object hierarchy
 - Server-Database-Schema-Object
- Permissions
 - DML: INSERT, UPDATE, DELETE,...
 - DDL: ALTER, ALTER ANY,...
 - Permission at different levels:
 - Server level: ALTER ANY DATABASE,...
 - Database level: BACKUP DATABASE,...
- Roles and Users
- GRANT, DENY and REVOKE



Physical Database Design

- Data Storage: Overview
- Indexes: Overview
- Query Execution Plans: Overview



Database Design Process revisit

- Database design process consists of the following main steps:
 - Requirements Analysis
 - Conceptual Database Design
 - Logical Database Design
 - Physical Database Design



Physical Database Design

- Physical Database Design pertains to developing an implementation for the database targeting a specified DBMS
- At this stage, the database designer decides on implementation being well-aware of the features, capabilities and limitations of the *target DBMS*.



Physical Database Design (contd.)

- This process includes:
 1. Translating logical data model for target DBMS (design base relations, constraints, derived data)
 2. Design file organisations and indexes focusing on performance
 3. Design user views and security measures

*Steps 1 and 3 has been covered to a certain extent in previous lectures/tutes/pracs

** Today, we'll introduce step 2



Design File Organizations and Indexes

- Determine and implement the optimal file organizations to store the base relations and the indexes that are required to achieve acceptable **performance**;
- Must understand the typical *workload* that database must support.
- Workload includes:
 - Most frequent/important queries and updates
 - Performance requirement for them




Data Storage: Disks and Files

- How is data stored in databases?
 - Secondary storage (typically on disks)
- Data are brought into memory from disk prior to db operations
- In terms of performance, most significant is the time pertaining to disk read and write (disk I/Os)
- Most DBMSs focus on reducing the time and access to the disk by implementing varied techniques

Disks

- Disk has platters
- Platters contains many tracks
- Each track is divided into smaller sector where data is stored
- Number of sectors make a block which is read/written from disk to memory (read) and vice-versa (write)





Accessing a Disk Page

- Time to access (read/write) a disk block:
 - *seek time* (moving arms to position disk head on track)
 - *rotational delay* (waiting for block to rotate under head)
 - *transfer time* (actually moving data to/from disk surface)

Disk I/O takes significantly more time than a memory I/O



Why Not Store Everything in Main Memory?

- *Costs too much.*
- *Main memory is volatile.* We want data to be saved between runs. (Obviously!)
- Typical storage hierarchy:
 - Main memory (RAM) for currently used data (primary storage).
 - Disk for the main database (secondary storage).
 - Disk for archiving older versions of the data (tertiary storage).




File Organization

- DBMS store tables as logical files
- Each *file* contains rows of the table
- Physically, a file contains a number of *pages*
- Size of page is same as size of block.

File Organization (contd.)

Example





File Organization

- Types of file organizations:
 - **Sequential file:** Order rows in the table according to certain attribute(s) values.
 - **Heap file:** Store rows in no particular order
 - **Hashed file:** Store rows based on a hash attribute(s) values




Indexes

- Indexes are auxiliary files that speed up selections on the rows of a table.
- Any subset of the fields of a relation can be the search key for an index on the relation.

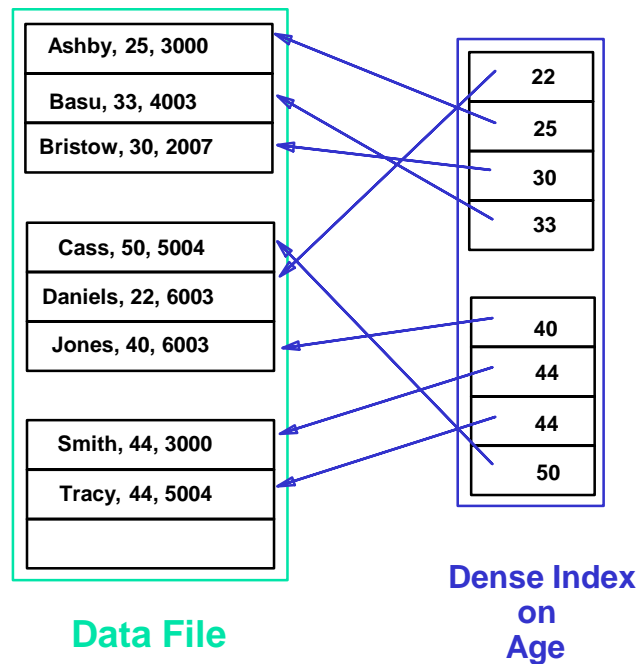
Indexes (contd.)

- Index file is smaller than the data file!
- Index file is always sorted or hashed!
- Searching the index using a search key field is usually efficient than searching a data file



Indexes (contd.)

■ Example





Indexes (contd.)

- Different types of data structures exist to build indexes which allow very efficient access to rows in a table
 - Example of data structures:
 - B+ Trees
 - Hash
 - Bitmapetc.
- * Discussing these in detail is beyond the scope of this course

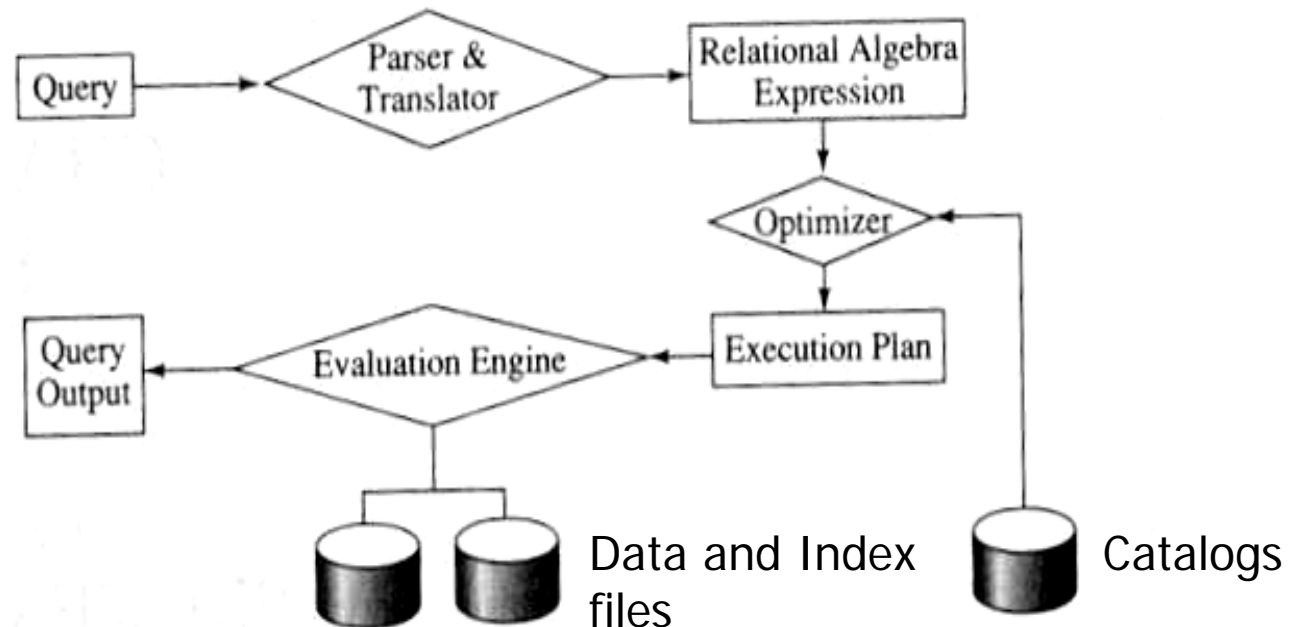


Query Execution Plans

- Every DBMS goes through a process called **query processing** prior to returning results when an SQL statement is given.

Steps in query processing

- The steps involved in query processing include...
 - Parsing and translation
 - Optimization
 - Evaluation





Steps in query processing (contd.)

- Step 1: Parsing and translation
 - An SQL statement is verified for correct syntax and semantics
 - The SQL statement is converted to a extended relational algebra expression



Steps in query processing (contd.)

- Step 2: Optimization
 - An efficient query execution strategy is devised (based on existing indexes, file organizations etc.)



Optimization

- Heart of performance in a DBMSs
- Involves complex algorithms and techniques
- A simple example:
 - Asking a balance of a particular account in a million page Account table, with appropriate indexes and a optimized plan, returns results within milliseconds while an hour in an unoptimized plan.



Steps in query processing (contd.)

- Step 3: Evaluation
 - The selected query execution plan is evaluated to obtain the results



Summary

- Physical Database Design pertaining to performance requires:
 - Analysis of the typical workload
 - An understanding of capabilities and limitations of target DBMS
 - A good understanding of storage structures, files organization, indexes and query execution plans



Lab this week

- Review and practice on Discretionary Access Control, and Physical Database Design



Issues and discussions on Assignment 3

- A3 specification
- A3 marking guide