

# Application Layer Protocols -1

---

A/PROF. DUY NGO

# Learning Objectives

---

## **2.1 Principles of network applications**

2.2 Web and HTTP

2.4 DNS

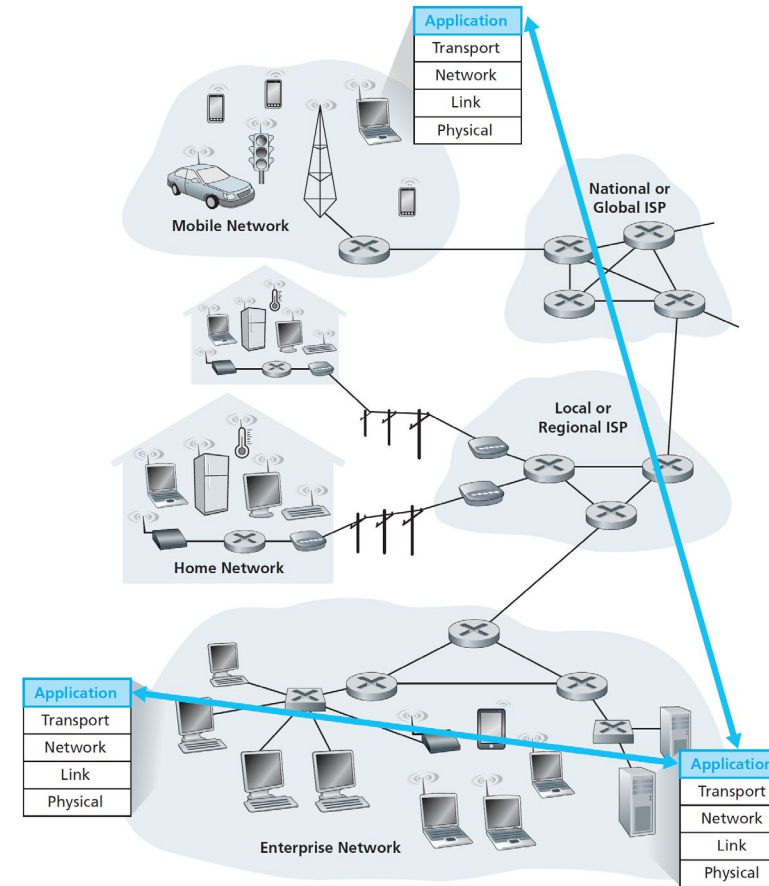
# Creating a Network App

**write programs that:**

- run on (different) **end systems**
- communicate over network
- e.g., web server software communicates with browser software

**no need to write software for network-core devices**

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



# Application Architectures

---

## **Possible structure of applications:**

- client-server
- peer-to-peer (P2P)


# Client-Server Architecture

## server:

- always-on host
- permanent IP address
- data centers for scaling


## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other



# P2P (Peer to Peer) Architecture

- **Minimum** reliance on dedicated servers
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - **self scalability** – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management



# Processes Communicating

---

**process:** program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS: Operating System)
- processes in different hosts communicate by exchanging **messages**

clients, servers


**client process:** process that initiates communication

**server process:** process that waits to be contacted

- aside: applications with P2P architectures have client processes & server processes

# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process sends message to out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process





# Addressing Processes

---

- to receive messages, process must have **identifier**
- host device has unique 32-bit IP address
- **Q:** does IP address of host on which process runs suffice for identifying the process?
  - **A:** no, **many** processes can be running on same host
- **identifier** includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - **IP address:** 128.119.245.12
  - **port number:** 80
- more shortly...

# App-layer protocol defines

---

- **types of messages exchanged,**
  - e.g., request, response
- **message syntax:**
  - what fields in messages & how fields are delineated
- **message semantics**
  - meaning of information in fields
- **rules** for when and how processes send & respond to messages

## **open protocols:**

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

## **proprietary protocols:**

- e.g., Skype

# What Transport Service Does An App Need?

---

## **data integrity**

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

## **timing**

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

## **throughput**

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

## **security**

- encryption, data integrity, ...

# Transport Service Requirements: Common Apps

| Application           | Data Loss     | Throughput                               | Time-Sensitive  |
|-----------------------|---------------|--|-----------------|
| file transfer         | no loss       | elastic                                  | no              |
| e-mail                | no loss       | elastic                                  | no              |
| Web documents         | no loss       | elastic                                  | no              |
| real-time audio/video | loss-tolerant | audio: 5kbps-1Mbps<br>video:10kbps-5Mbps | yes, 100's msec |
| stored audio/video    | loss-tolerant | same as above                            | yes, few secs   |
| interactive games     | loss-tolerant | few kbps up                              | yes, 100's msec |
| text messaging        | no loss       | elastic                                  | yes and no      |

# Internet Transport Protocols Services

---

## TCP service:

- **reliable transport** between sending and receiving process
- **flow control:** sender won't overwhelm receiver
- **congestion control:** throttle sender when network overloaded
- **does not provide:** timing, minimum throughput guarantee, security

- **connection-oriented:** setup required between client and server processes

## UDP service:

- **unreliable data transfer** between sending and receiving process
- **does not provide:** reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

**Q:** why bother? Why is there a UDP?

# Internet Apps: Application, Transport Protocols

---

| <b>Application</b>     | <b>Application layer protocol</b>    | <b>Underlying transport protocol</b> |
|------------------------|--------------------------------------|--------------------------------------|
| e-mail                 | SMTP [RFC 2821]                      | TCP                                  |
| remote terminal access | Telnet [RFC 854]                     | TCP                                  |
| Web                    | HTTP [RFC 2616]                      | TCP                                  |
| file transfer          | FTP [RFC 959]                        | TCP                                  |
| streaming multimedia   | HTTP (e.g., YouTube), RTP [RFC 1889] | TCP or UDP                           |
| Internet telephony     | SIP, RTP, proprietary (e.g., Skype)  | TCP or UDP                           |

# Securing TCP

---

## TCP & UDP

- no encryption
- cleartext passwds sent into socket traverse Internet in cleartext

## SSL

- provides encrypted TCP connection
- data integrity
- end-point authentication

## SSL is at app layer

- apps use SSL libraries, that “talk” to TCP

## SSL socket API

- cleartext passwords sent into socket traverse Internet encrypted
- see Chapter 8

# Web and HTTP

---

## First, a review...

- **web page** consists of **objects**
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of **base HTML-file** which includes **several referenced objects**
- each object is addressable by a **URL**, e.g.,

www.someschool.edu / someDept/pic.gif  
host name                      path name




# HTTP Overview (1 of 2)

## HTTP: hypertext transfer protocol

Web's application layer protocol

client/server model

- **client:** browser that requests, receives, (using HTTP protocol) and “displays” Web objects
- **server:** Web server sends (using HTTP protocol) objects in response to requests



# HTTP Overview (2 of 2)

---

## uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

## HTTP is “stateless”

- server maintains no information about past client requests

## aside

### protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

# HTTP Connections

---

## non-persistent HTTP

- at most one object sent over TCP connection
  - connection then closed
- downloading multiple objects required multiple connections

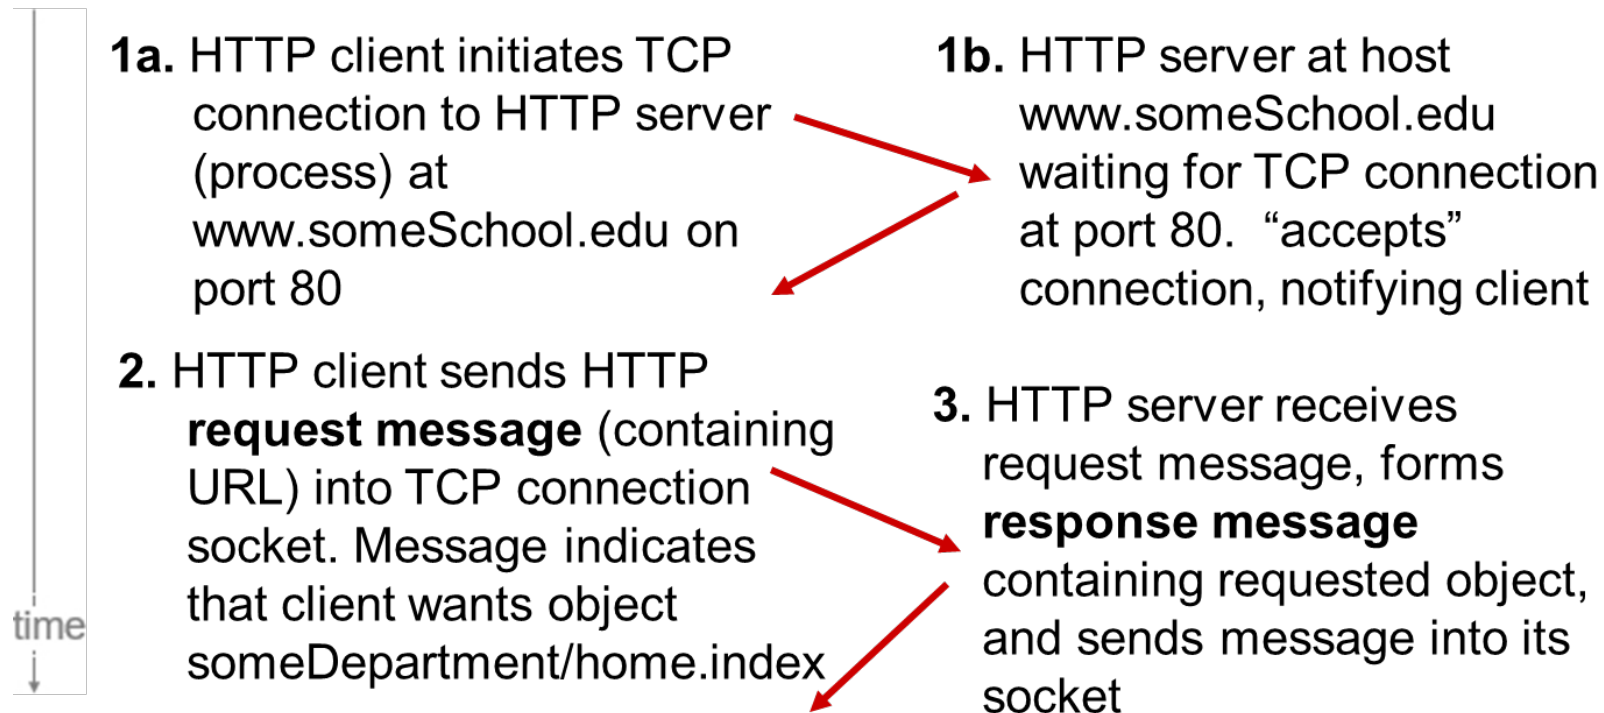
## persistent HTTP

- multiple objects can be sent over single TCP connection between client, server

# Non-Persistent HTTP (1 of 2)

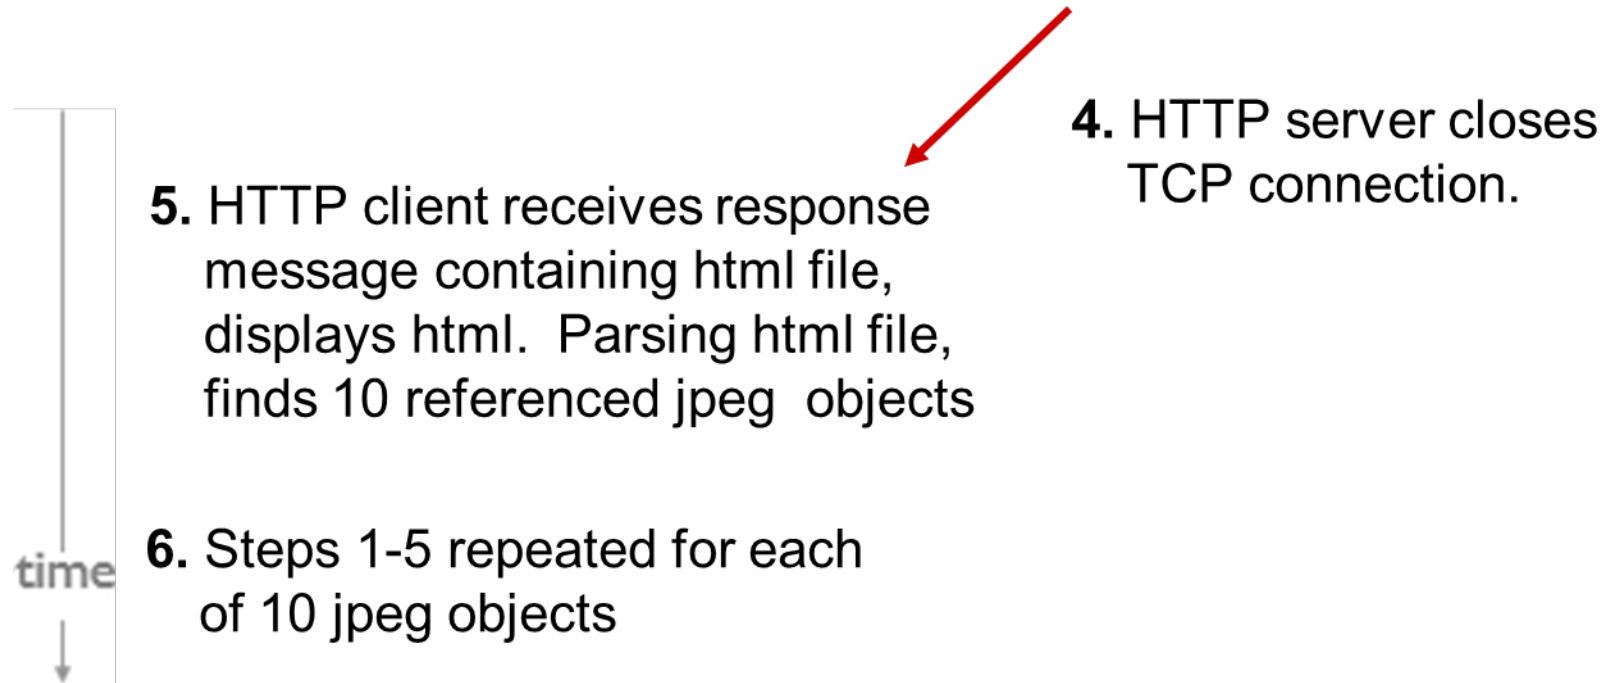
suppose user enters URL:

[www.someSchool.edu/someDepartment/home.index](http://www.someSchool.edu/someDepartment/home.index)



# Non-Persistent HTTP (2 of 2)

---



# Non-Persistent HTTP: Response Time

**RTT (definition):** time for a small packet to travel from client to server and back


**HTTP response time:**

one RTT to initiate TCP connection

one RTT for HTTP request and first few bytes of HTTP response to return

file transmission time

non-persistent HTTP response time =  $2\text{RTT} + \text{file transmission time}$



# Persistent HTTP

---

## non-persistent HTTP issues:

- requires 2 RTT s per object
- OS overhead for **each** TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## persistent HTTP:

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

# HTTP Request Message

- two types of HTTP messages: **request**, **response**
- **HTTP request message:**
  - ASCII (human-readable format)

The diagram shows an example of an HTTP request message in ASCII format. Annotations with arrows point to specific parts of the message:

- request line (GET, POST, HEAD commands):** Points to the first line: `GET /index.html HTTP/1.1\r\n`
- header lines:** A bracket on the left groups the following lines: `Host: www-net.cs.umass.edu\r\n`, `User-Agent: Firefox/3.6.10\r\n`, `Accept: text/html,application/xhtml+xml\r\n`, `Accept-Language: en-us,en;q=0.5\r\n`, `Accept-Encoding: gzip,deflate\r\n`, `Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n`, `Keep-Alive: 115\r\n`, and `Connection: keep-alive\r\n`
- carriage return, line feed at start of line indicates end of header lines:** Points to the `\r\n` at the end of the `Connection` header line.
- carriage return character** and **line-feed character:** Arrows point to the `\r` and `\n` characters at the end of the first line.

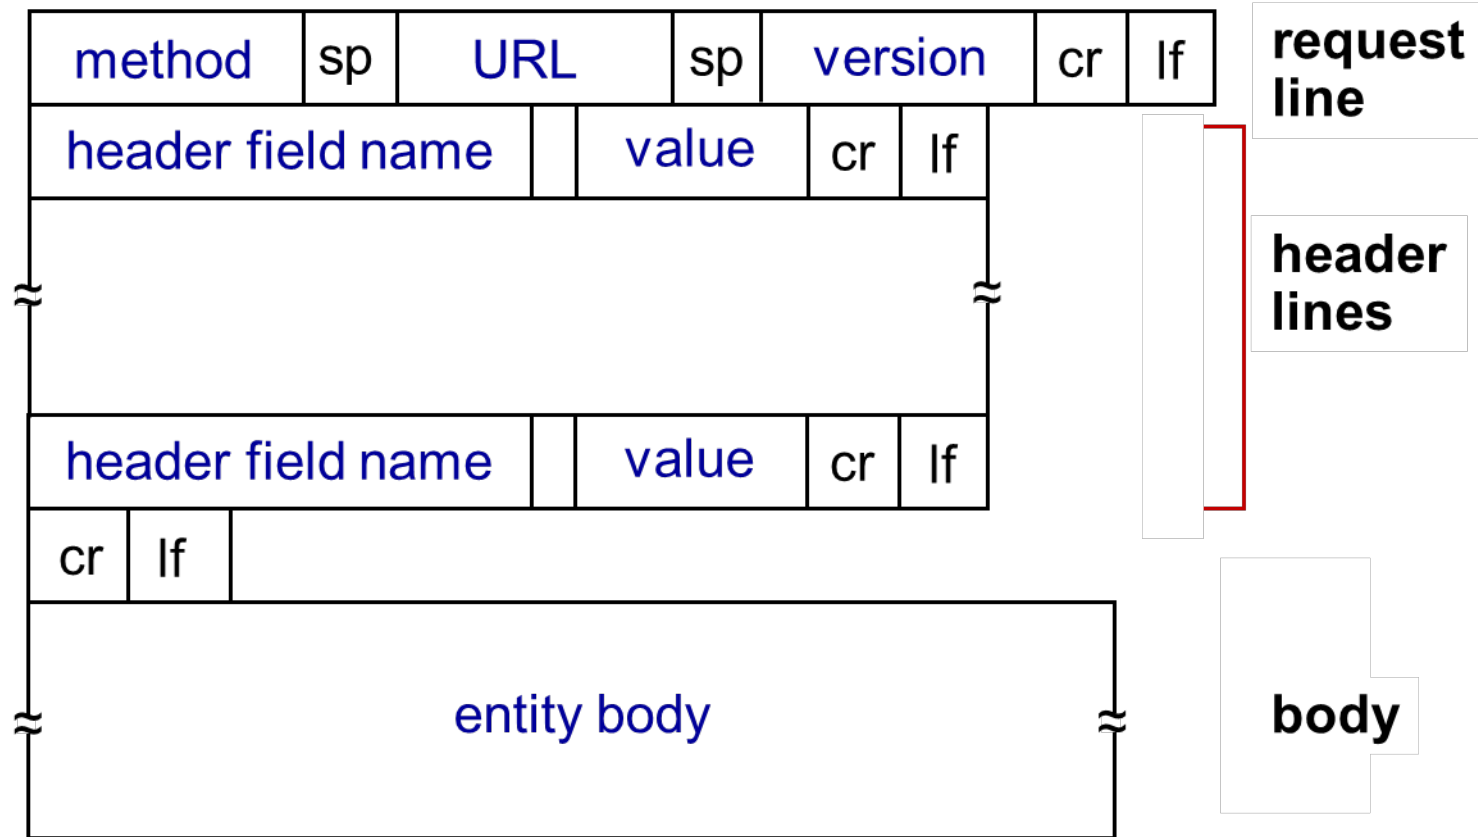
```
GET /index.html HTTP/1.1\r\nHost: www-net.cs.umass.edu\r\nUser-Agent: Firefox/3.6.10\r\nAccept: text/html,application/xhtml+xml\r\nAccept-Language: en-us,en;q=0.5\r\nAccept-Encoding: gzip,deflate\r\nAccept-Charset: ISO-8859-1,utf-8;q=0.7\r\nKeep-Alive: 115\r\nConnection: keep-alive\r\n\r\n
```

\* Check out the online interactive exercises for more examples:

[http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)



# HTTP Request Message: General Format



# DNS: Domain Name System

---

**people:** many identifiers:

- SSN, name, passport #

**Internet hosts, routers:**

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., [www.yahoo.com](http://www.yahoo.com) - used by humans

**Q:** how to map between IP address and name, and vice versa?

**Domain Name System:**

- **distributed database** implemented in hierarchy of many **name servers**
- **application-layer protocol:** hosts, name servers communicate to **resolve** names (address/name translation)
  - note: core Internet function, implemented as application-layer protocol
  - complexity at network’s “edge”

# DNS: Services, Structure

---

## DNS services

- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - replicated Web servers:  
many IP addresses  
correspond to one name


## why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

**A: doesn't scale!**

# DNS: A Distributed, Hierarchical Database

---



client wants IP for [www.amazon.com](http://www.amazon.com);

**1<sup>st</sup> approximation:**

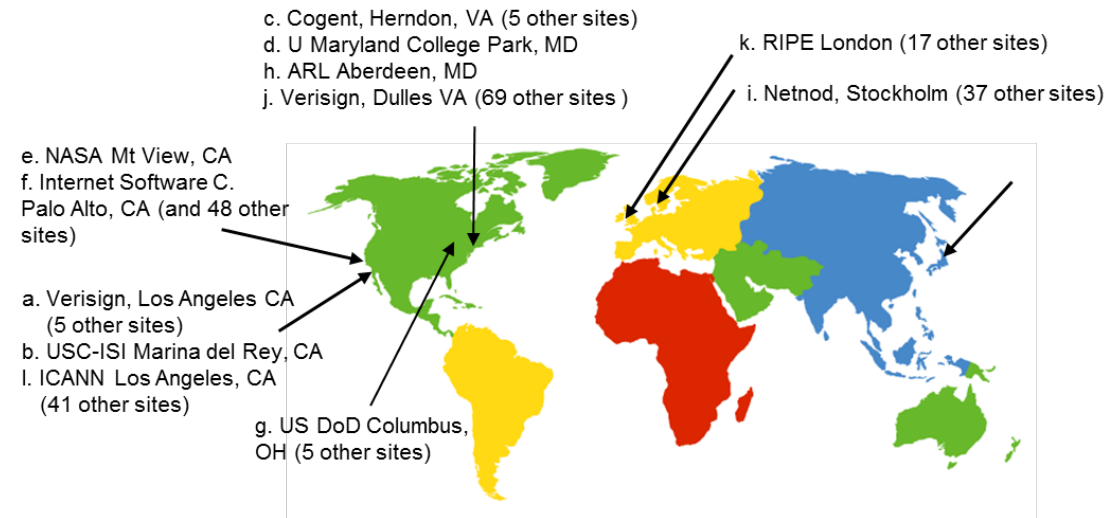
- client queries root server to find com DNS server
- client queries .com DNS server to get [amazon.com](http://amazon.com) DNS server
- client queries [amazon.com](http://amazon.com) DNS server to get IP address for [www.amazon.com](http://www.amazon.com)

# DNS: Root Name Servers

- contacted by local name server that can not resolve name
- root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server

## 13 logical root name “servers” worldwide

- each “server” replicated many times



# TLD, Authoritative Servers

---

## **top-level domain (TLD) servers:**

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

## **authoritative DNS servers:**

- Organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

# Local DNS Name Server

---

does not strictly belong to hierarchy

each ISP (residential ISP, company, university) has one

- also called “default name server”

when host makes DNS query, query is sent to its local DNS server


- has local cache of recent name-to-address translation pairs (but may be out of date!)
- acts as proxy, forwards query into hierarchy

# DNS Name Resolution Example (1 of 2)

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## iterated query:

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”






# DNS Name Resolution Example (2 of 2)

## recursive query:

puts burden of name resolution on contacted name server

heavy load at upper levels of hierarchy?



# DNS: Caching, Updating Records

---

once (any) name server learns mapping, it **caches** mapping

- cache entries timeout (disappear) after some time (TTL)
- TLD servers typically cached in local name servers
- thus root name servers not often visited

cached entries may be **out-of-date** (best effort name-to-address translation!)

- if name host changes IP address, may not be known Internet-wide until all TTL s expire

update/notify mechanisms proposed IETF standard

- RFC 2136

# DNS Records

---

**DNS:** distributed database storing resource records (**RR**)

RR format: (**name**, **value**, **type**, **ttl**)

## **type=A**

- **name** is hostname
- **value** is IP address

## **type=NS**

- **name** is domain (e.g., [foo.com](http://foo.com))
- **value** is hostname of authoritative name server for this domain

## **type=CNAME**

- **name** is alias name for some “canonical” (the real) name
- [www.ibm.com](http://www.ibm.com) is really [servereast.backup2.ibm.com](http://servereast.backup2.ibm.com)
- **value** is canonical name

## **type=MX**

- **value** is name of mailserver associated with **name**

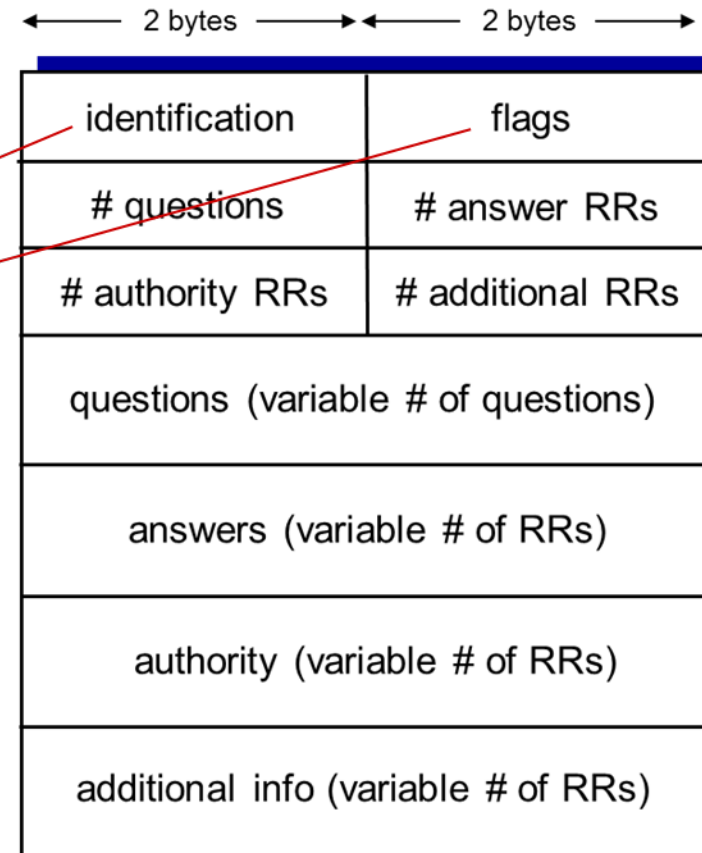
# DNS Protocol, Messages (1 of 2)

- **query** and **reply** messages, both with same **message format**

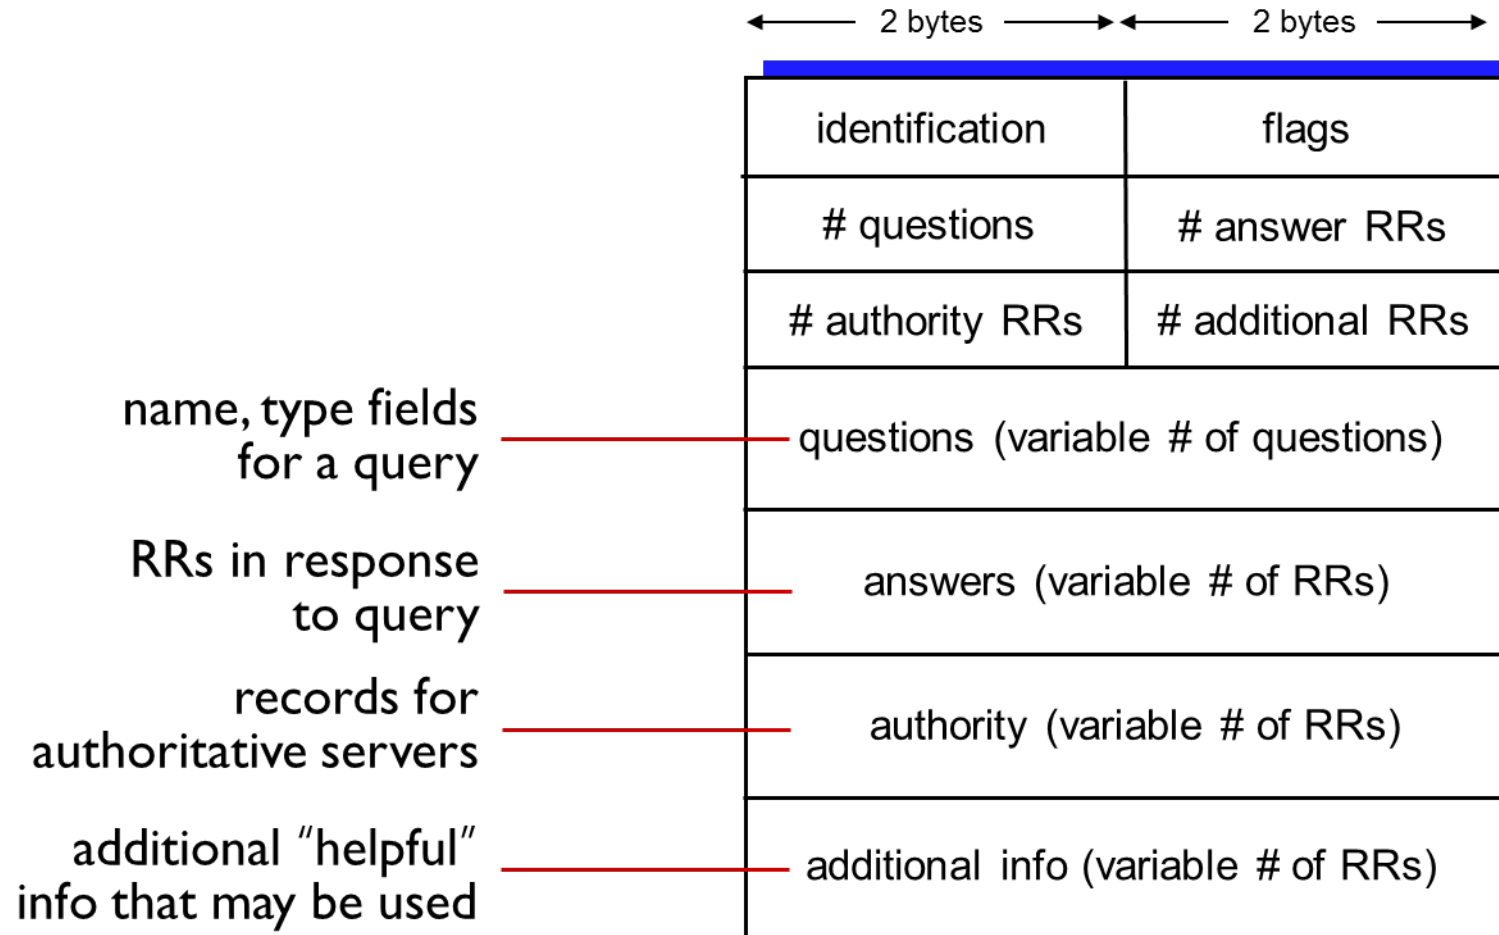
message header

- **identification:** 16 bit # for query, reply to query uses same #

- **flags:**
  - query or reply
  - recursion desired
  - recursion available
  - reply is authoritative



# DNS Protocol, Messages (2 of 2)



# Attacking DNS

---

## **DDoS attacks**

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

## **redirect attacks**

- man-in-middle
  - Intercept queries
- DNS poisoning
  - Send bogus replies to DNS server, which caches

## **exploit DNS for DDoS**

- send queries with spoofed source address: target IP
- requires amplification