

Network Layer: Routing Protocols

A/PROF. DUY NGO

Learning Objectives

5.1 introduction

5.2 routing protocols

- link state
- distance vector

Network-Layer Functions

Recall: two network-layer functions:

- **forwarding:** move packets from router's input to appropriate router output **data plane**


routing: determine route taken by packets from source to destination **control plane**

Two approaches to structuring network control plane:

- per-router control (traditional)
- logically centralized control (software defined networking)


Per-Router Control Plane

Individual routing algorithm components **in each and every router** interact with each other in control plane to compute forwarding tables



Logically Centralized Control Plane

A distinct (typically remote) controller interacts with local control agents (CAs) in routers to compute forwarding tables




Routing Protocols

Routing protocol goal: determine “good” paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- path: sequence of routers packets will traverse in going from given initial source host to given final destination host
- “good”: least “cost”, “fastest”, “least congested”
- routing: a “top-10” networking challenge!

Graph Abstraction of the Network




graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

aside: graph abstraction is useful in other network contexts, e.g., P2P, where \mathbf{N} is set of peers and \mathbf{E} is set of TCP connections

Graph Abstraction: Costs



$c(x, x') = \text{cost of link } (x, x')$ e.g., $c(w, z) = 5$

cost could always be 1, or
inversely related to bandwidth,
or inversely related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z?

routing algorithm: algorithm that finds that least cost path

Routing Algorithm Classification

Q: global or decentralized information?

global:

- all routers have complete topology, link cost info
- **“link state” algorithms**

decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- **“distance vector” algorithms**

Q: static or dynamic?

static:

- routes change slowly over time

dynamic:

- routes change more quickly
 - periodic update
 - in response to link cost changes

A Link-State Routing Algorithm

Dijkstra's algorithm

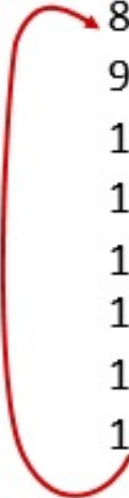
- network topology, link costs: known to all nodes
 - accomplished via “link state broadcast”
 - all nodes have the same information
- compute the least-cost paths from one node (“source”) to all other nodes
 - gives **forwarding table** for that node
- iterative: after k iterations, know the least-cost paths to k destinations

Notation:

- **$c(x,y)$** : link cost between node x and node y
- If **(x,y)** does not belong to E (that is, x and y are not direct neighbours), then we set **$c(x,y) = \infty$**
- **$D(v)$** : current cost of the least-cost path from the source to destination v (as of this iteration of the routing algorithm)
- **$p(v)$** : previous node (neighbour of v) along the current least-cost path from the source to v
- **N'** : a subset of nodes whose least-cost path is definitively known

Dijkstra's Algorithm

```
1 Initialization:
2   $N' = \{u\}$ 
3  for all nodes  $v$ 
4    if  $v$  adjacent to  $u$ 
5      then  $D(v) = c(u,v)$ 
6    else  $D(v) = \infty$ 
7
8  Loop
9    find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10   add  $w$  to  $N'$ 
11   update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$  :
12      $D(v) = \min( D(v), D(w) + c(w,v) )$ 
13   /* new cost to  $v$  is either old cost to  $v$  or known
14   shortest path cost to  $w$  plus cost from  $w$  to  $v$  */
15 until all nodes in  $N'$  (that is,  $N'=N$ )
```




Dijkstra's Algorithm: Example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	<u>3,u</u>	5,u	∞	∞
1	uw	6,w		<u>5,u</u>	11,w	∞
2	uwX	<u>6,w</u>			11,w	14,x
3	uwxv				<u>10,y</u>	14,x
4	uwxvy					<u>12,y</u>
5	uwxvyz					


Notes:

- construct shortest path tree by tracing predecessor nodes
- ties can exist (can be broken arbitrarily)



Dijkstra's Algorithm: Another Example


Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux ←	2,u	4,x		2,x	∞
2	uxy ←	2,u	3,y			4,y
3	uxyv ←		3,y			4,y
4	uxyvw ←					4,y
5	uxyvwz ←					



* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

Dijkstra's Algorithm: Example (2 of 2)

- The resulting **least-cost paths** from u:



- The resulting **forwarding table** in u:

Destination	Link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

Dijkstra's Algorithm: Discussion

Computational complexity: given n nodes (not counting the source)

- First iteration: need to check all n nodes to find the node w not in N' that has the minimum cost.
- Second iteration: check $(n-1)$ nodes to determine the minimum cost
- And so on...
- Overall, $n(n+1)/2$ comparisons. Worst-case complexity: $O(n^2)$
- More efficient implementation (using heap data structure) to reduce complexity to $O(n\log(n))$

Distance Vector Algorithm (1 of 6)

Bellman-Ford equation (dynamic programming)

let

$d_x(y) :=$ cost of least-cost path from x to y

then


$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

\min taken over all neighbors v of x

cost to neighbor v

cost from neighbor v to destination y

Bellman-Ford Example



clearly, $d_v(z)=5$, $d_x(z)=3$, $d_w(z)=3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node achieving minimum is the next hop in the shortest path,
used in the forwarding table

Distance Vector Algorithm (2 of 6)

Initialisation step:

- Each node x begins with $D_x(y)$, an estimate of the cost of the least-cost path from x to y , for all nodes y in N .
- Let $\mathbf{D}_x = [D_x(y): y \text{ in } N]$ be node x 's distance vector of cost estimates from x to all other nodes y in N .
- Each node x maintains the following routing information:
 - The cost $c(x,v)$ from x to each directly attached neighbor v
 - Node x 's distance vector $\mathbf{D}_x = [D_x(y): y \text{ in } N]$
 - The distance vectors of each of its neighbours v , that is, $\mathbf{D}_v = [D_v(y): y \text{ in } N]$

Distance Vector Algorithm (3 of 6)

Update step:

- From time-to-time, each node sends a copy of its own distance vector to its neighbours
- When a node x receives a new distance vector from its neighbor w , it saves w 's distance vector and updates its own distance vector using B-F equation: $D_x(y) \leftarrow \min_v \{c(x, v) + D_v(y)\}$ for each node $y \in N$
- If node x 's distance vector has changed as a result of this update step, node x sends its updated distance vector to each of its neighbours, which in turn update their own distance vectors.
- Miraculously, each cost **estimate** $D_x(y)$ **converges** to $d_x(y)$, the **actual** cost of the least-cost path from node x to node y !

Distance Vector Algorithm (4 of 6)

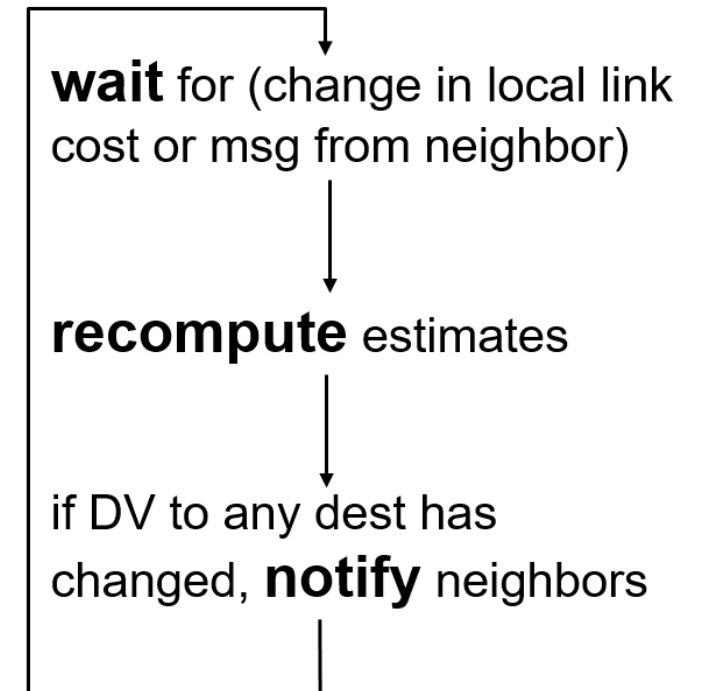
iterative, asynchronous: each local iteration caused by:

- local link cost change
- Distance vector update message from neighbour


distributed:

- each node notifies neighbors **only** when its distance vector changes
 - neighbors then notify their neighbors if necessary

each node:



Distance Vector Algorithm (5 of 6)



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\} \\ = \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\} \\ = \min\{2+1, 7+0\} = 3$$

node x table

	cost to		
	x	y	z
from x	0	2	7
from y	∞	∞	∞
from z	∞	∞	∞

node x table (updated)

	cost to		
	x	y	z
from x	0	2	3
from y	2	0	1
from z	7	1	0

node y table



	cost to		
	x	y	z
from x	∞	∞	∞
from y	2	0	1
from z	∞	∞	∞

node z table

	cost to		
	x	y	z
from x	∞	∞	∞
from y	∞	∞	∞
from z	7	1	0

time →


Distance Vector Algorithm (6 of 6)



Distance Vector: Link Cost Changes (1 of 2)

link cost changes:

- node detects local link cost change
- updates routing info, recalculates distance vector
- if DV changes, notify neighbors



“good news travels fast”

t_0 : y detects link-cost change, updates its DV, informs its neighbors.

t_1 : z receives update from y, updates its table, computes new least cost to x, sends its neighbors its DV.

t_2 : y receives z's update, updates its distance table. y's least costs do *not* change, so y does **not** send a message to z.


* Check out the online interactive exercises for more examples:

http://gaia.cs.umass.edu/kurose_ross/interactive/

Distance Vector: Link Cost Changes (2 of 2)

link cost changes:

- node detects local link cost change
- **bad news travels slow** – “count to infinity” problem!
- 44 iterations before algorithm stabilizes: see text



poisoned reverse:

- If z routes through y to get to x:
 - z tells y its (z's) distance to x is infinite (so y won't route to x via z)
- Will this completely solve count to infinity problem? No – for loops involving three or more nodes.

Comparison of LS and DV Algorithms

message complexity

- **LS:** with n nodes, E links, $O(nE)$ msgs sent
- **DV:** exchange between neighbors only
 - convergence time varies

speed of convergence

- **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect **link** cost
- each node computes only its **own** table

DV:

- DV node can advertise incorrect **path** cost
- each node's table used by others
 - error propagate thru network